

CAGE, ggplot, and Tidyverse

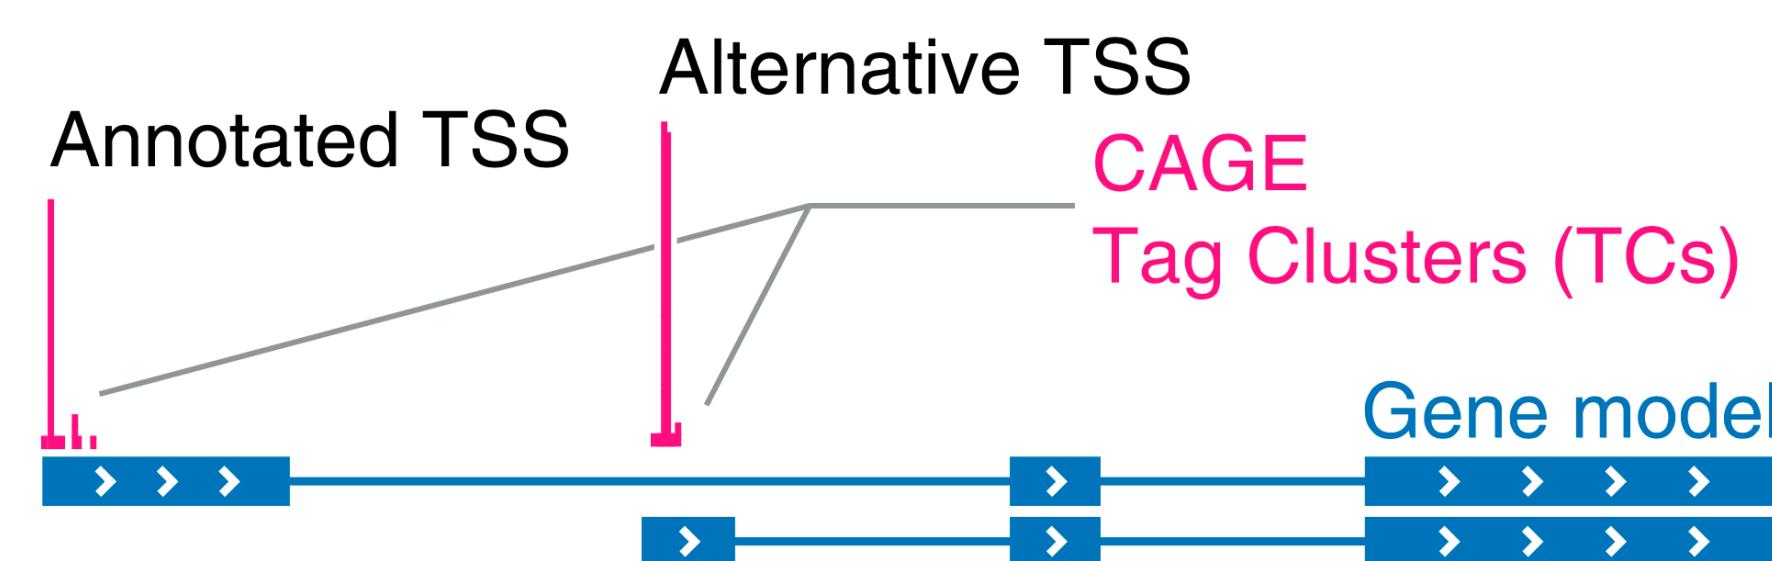
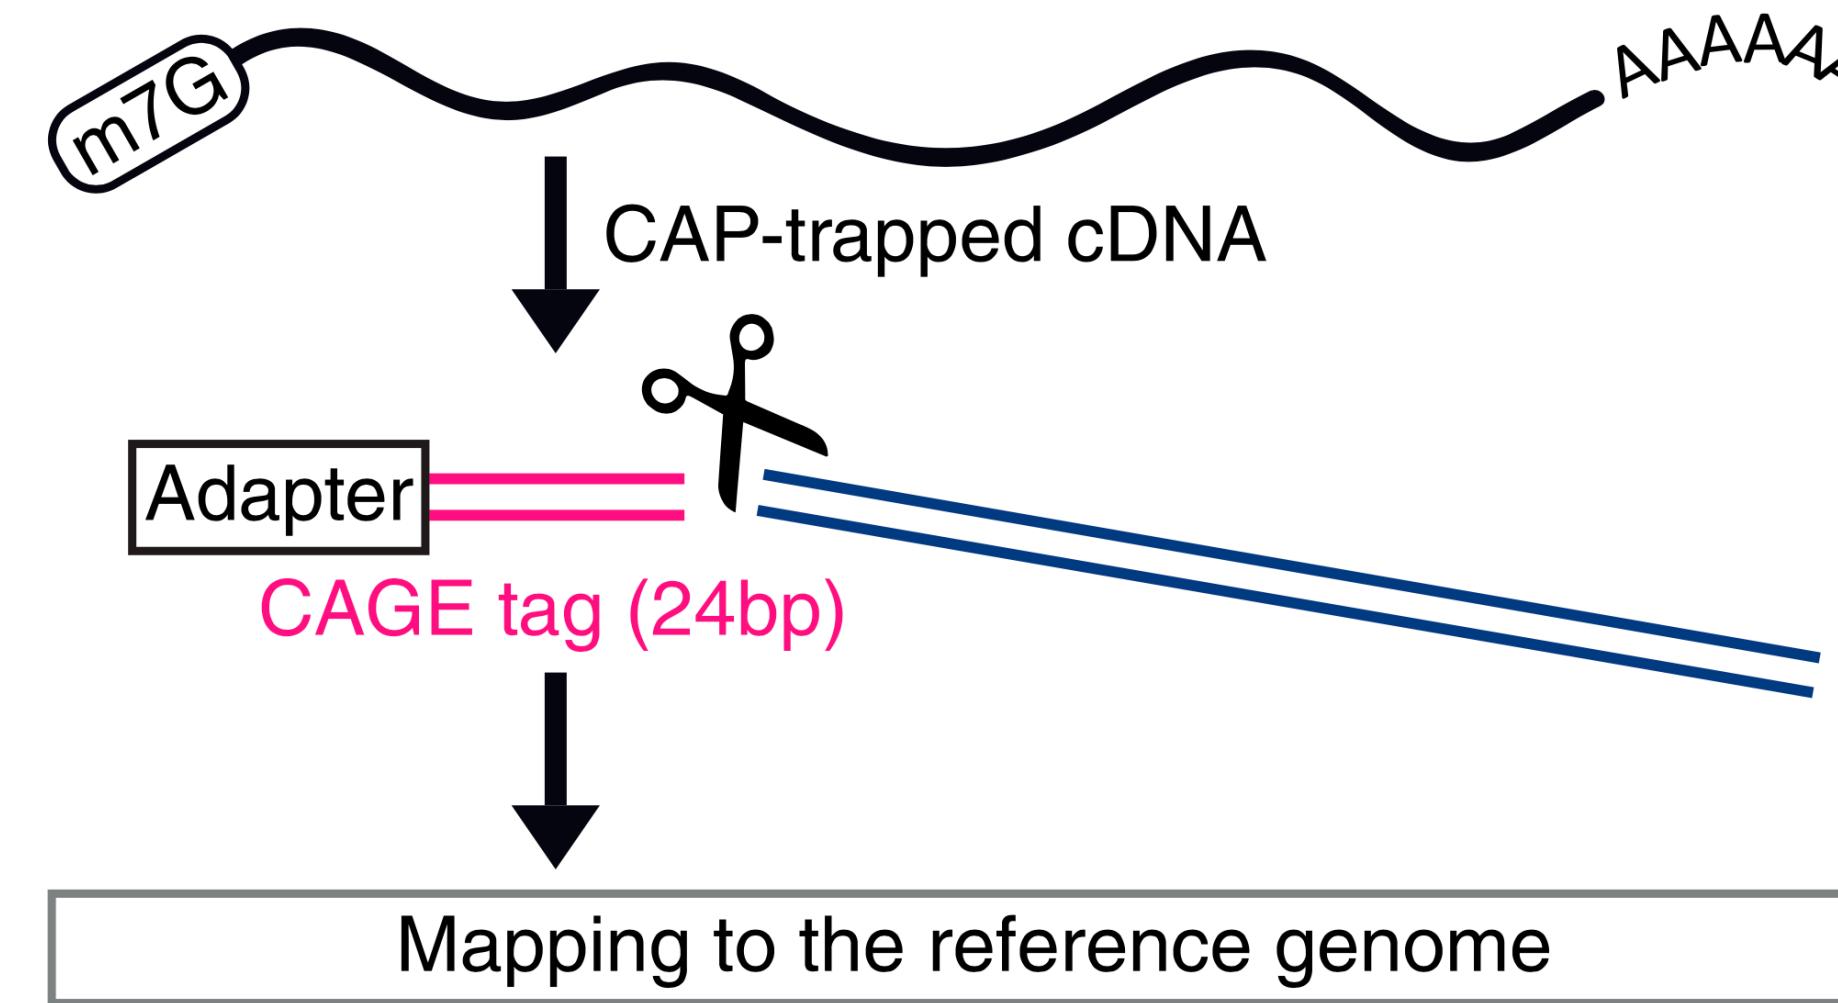
Axel Thieffry, PhD

October 2021

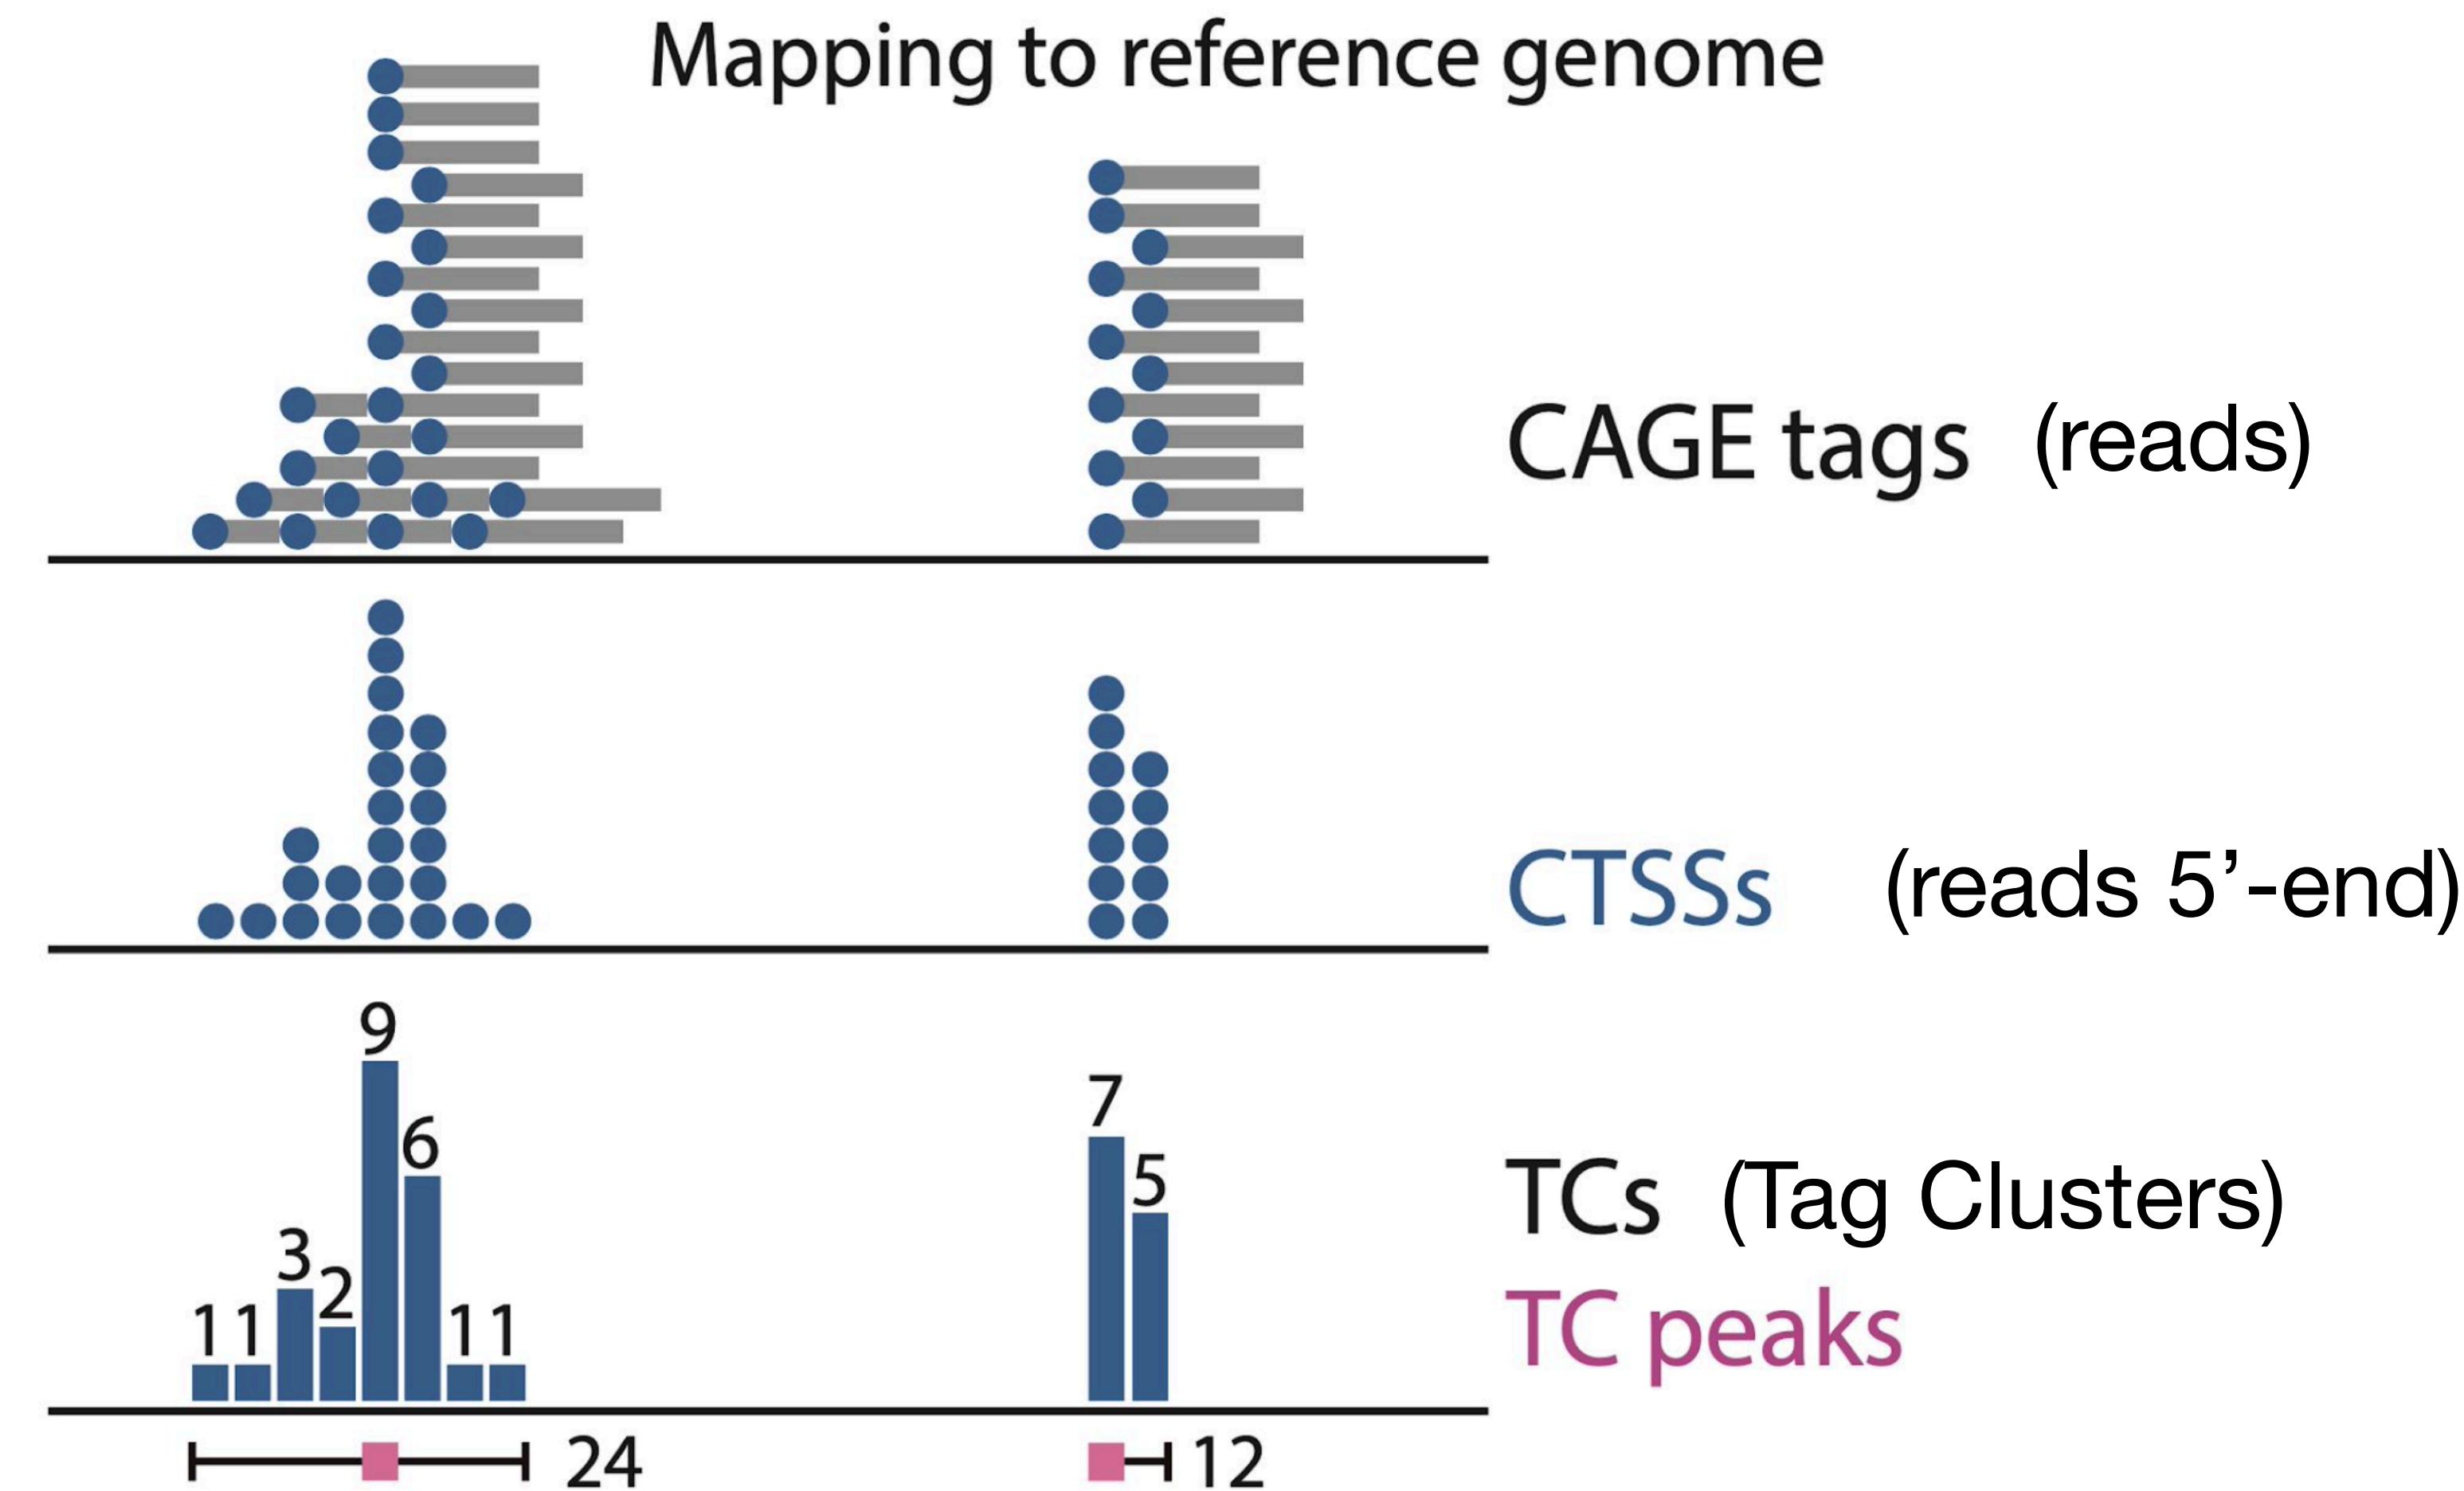
CAGE intro

Cap Analysis of Gene Expression

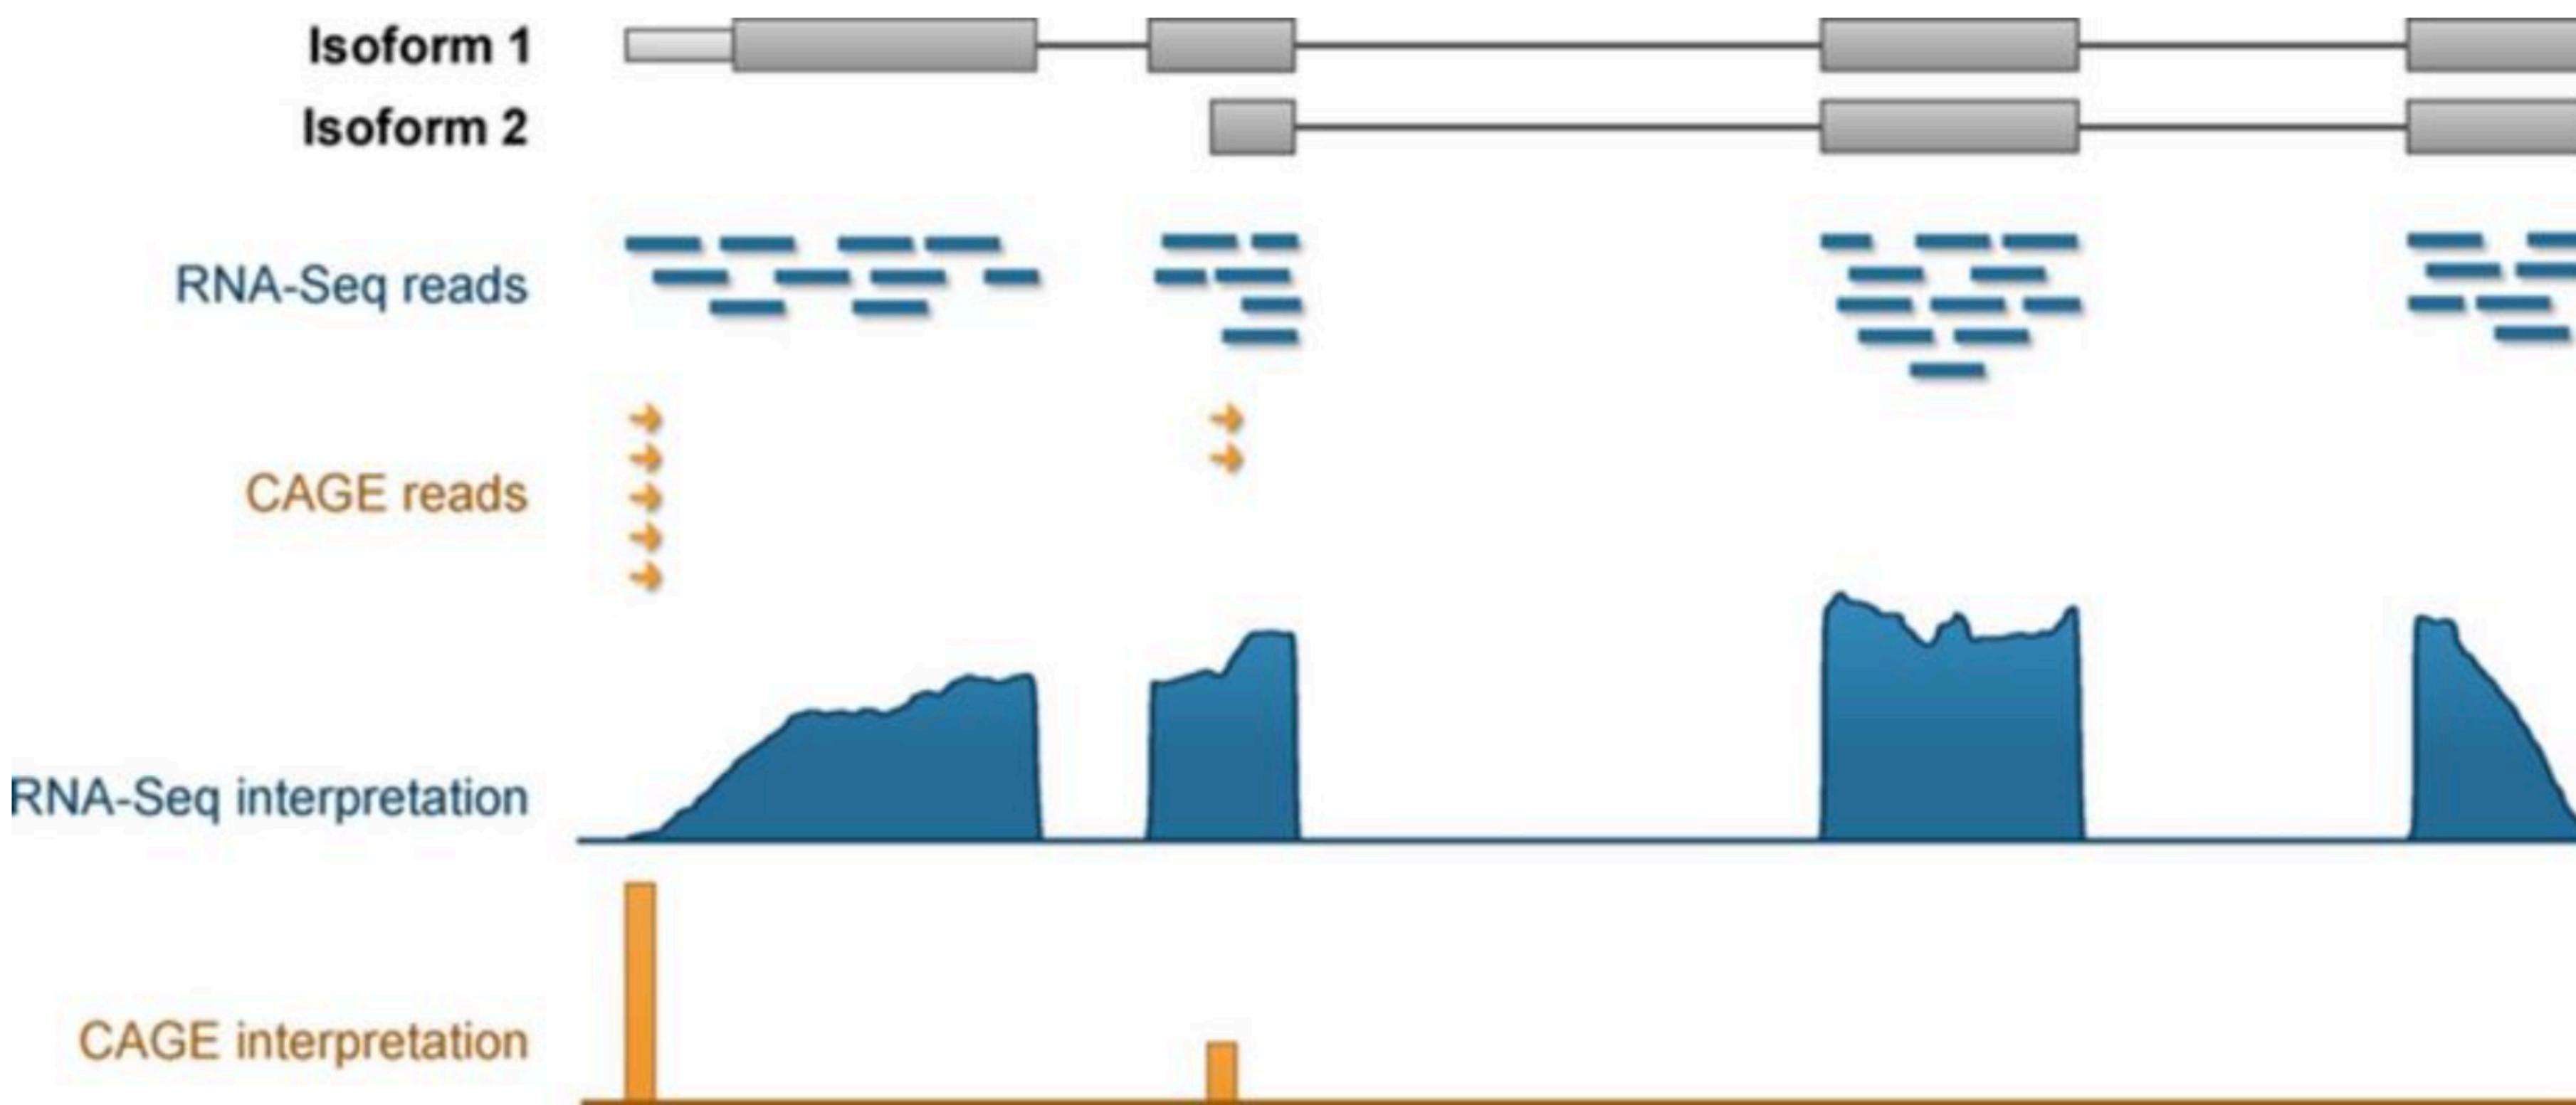
5'-end sequencing



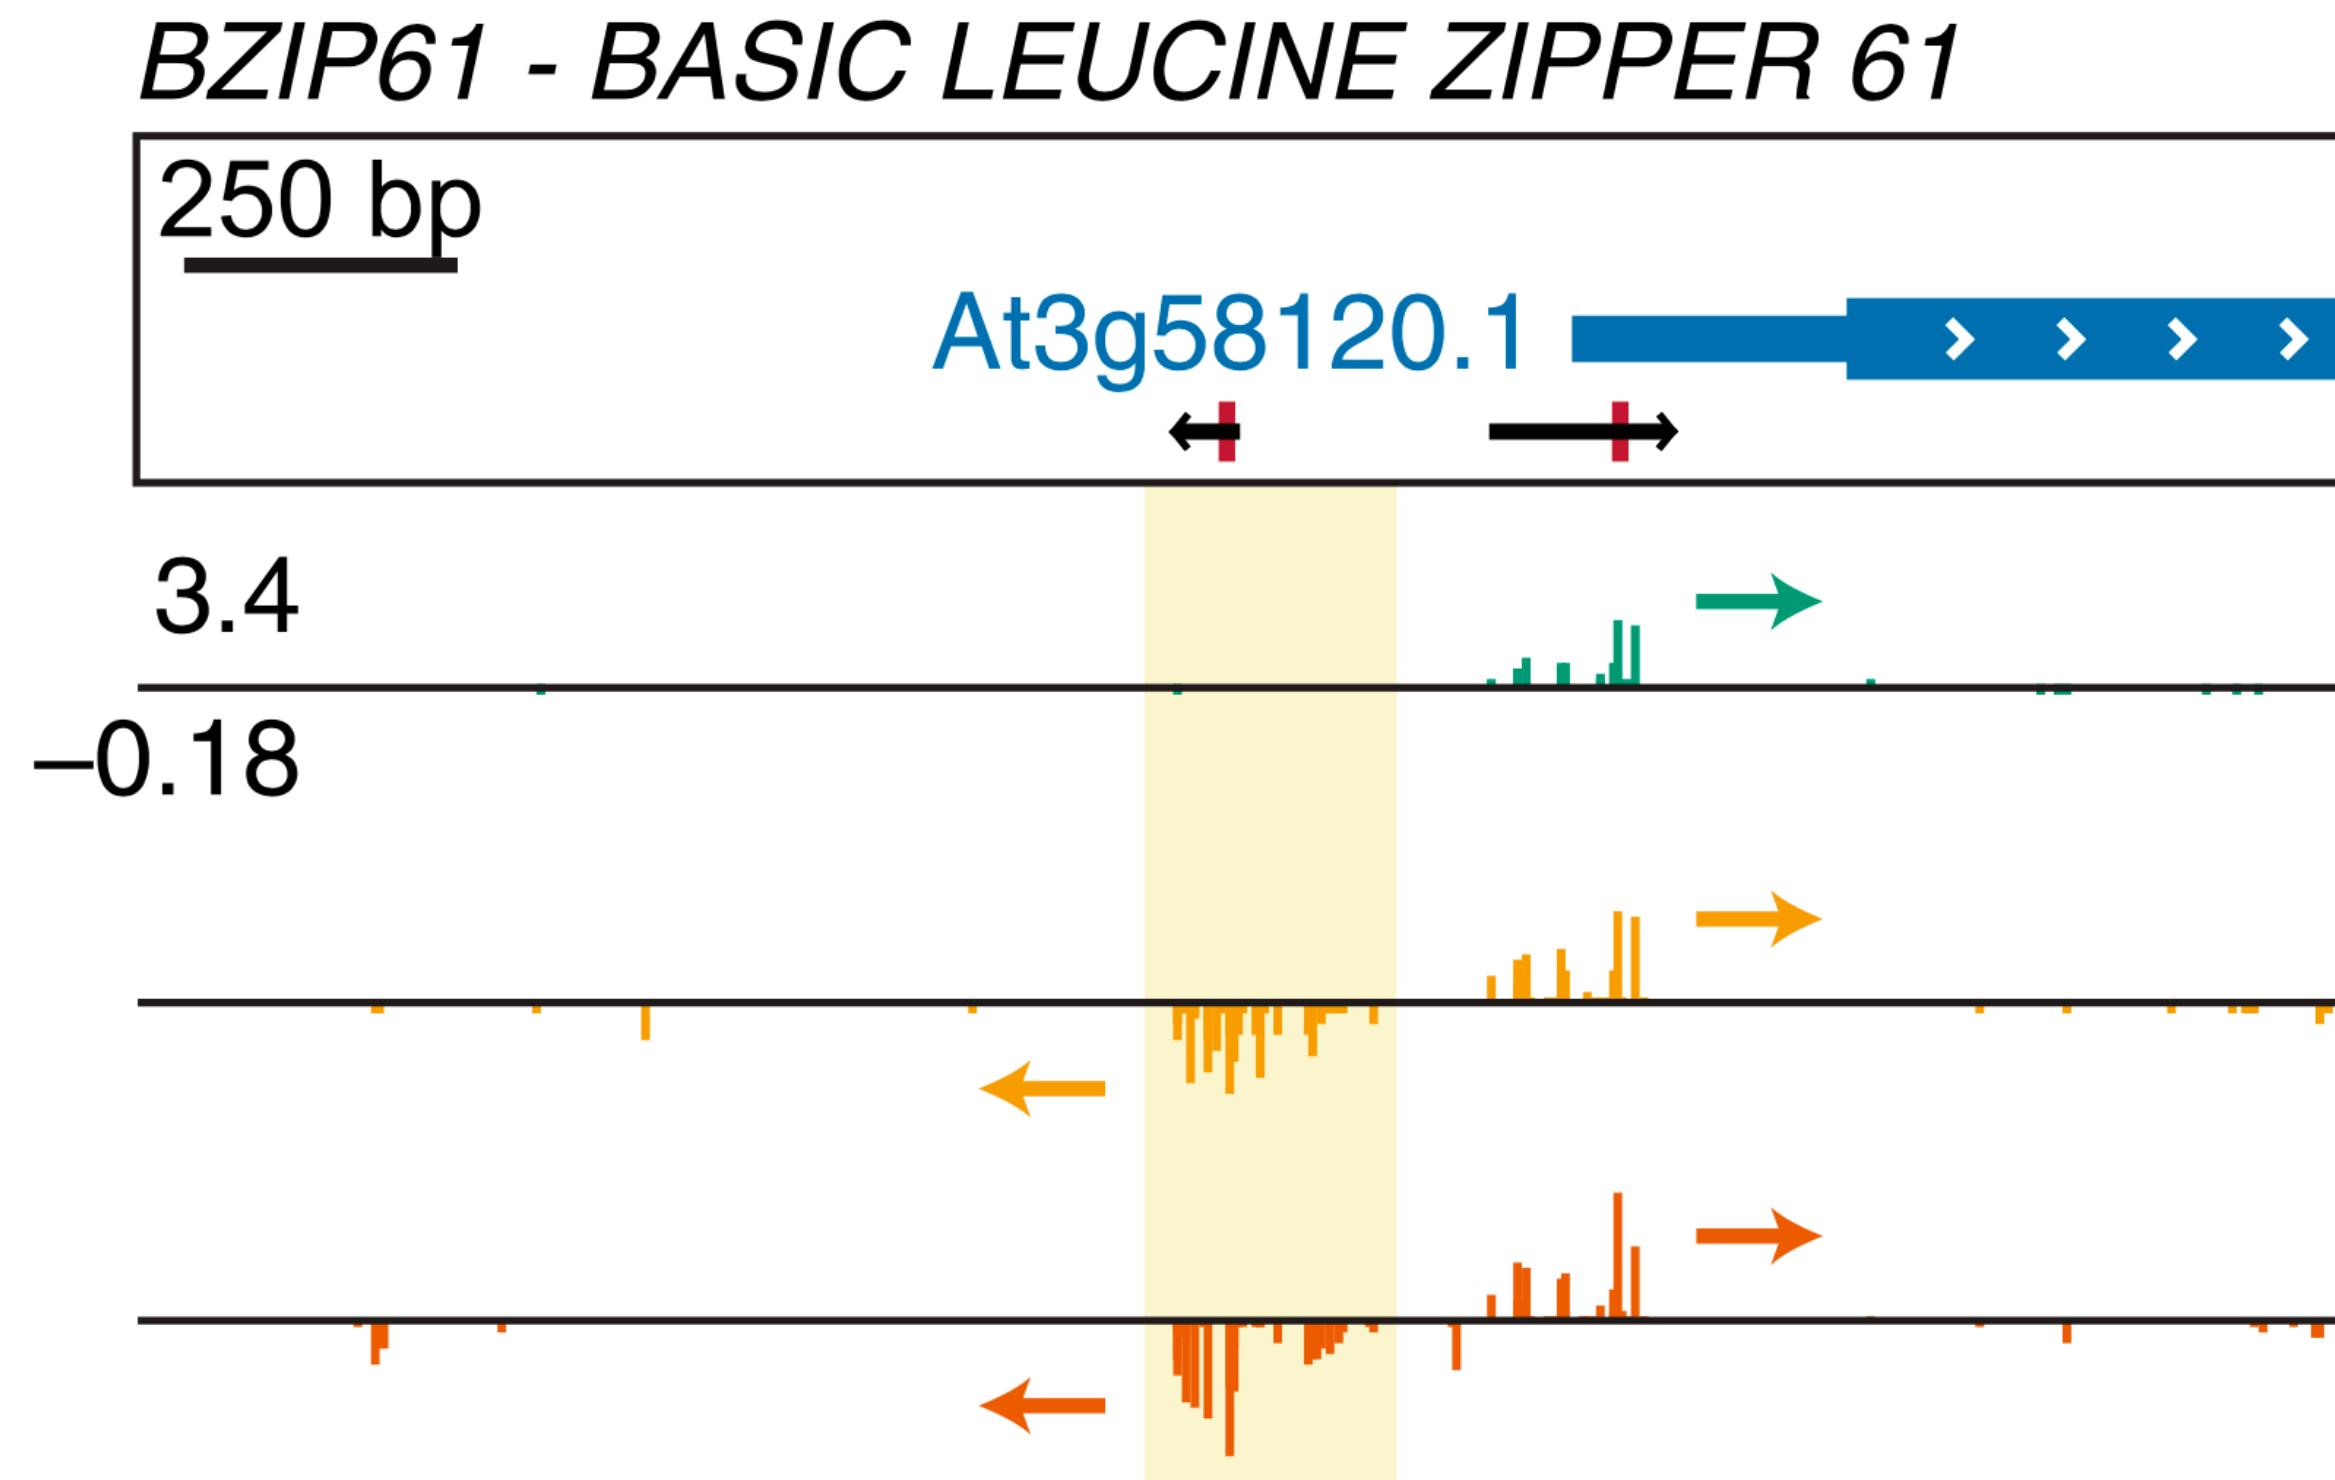
Conceptual CAGE pipeline



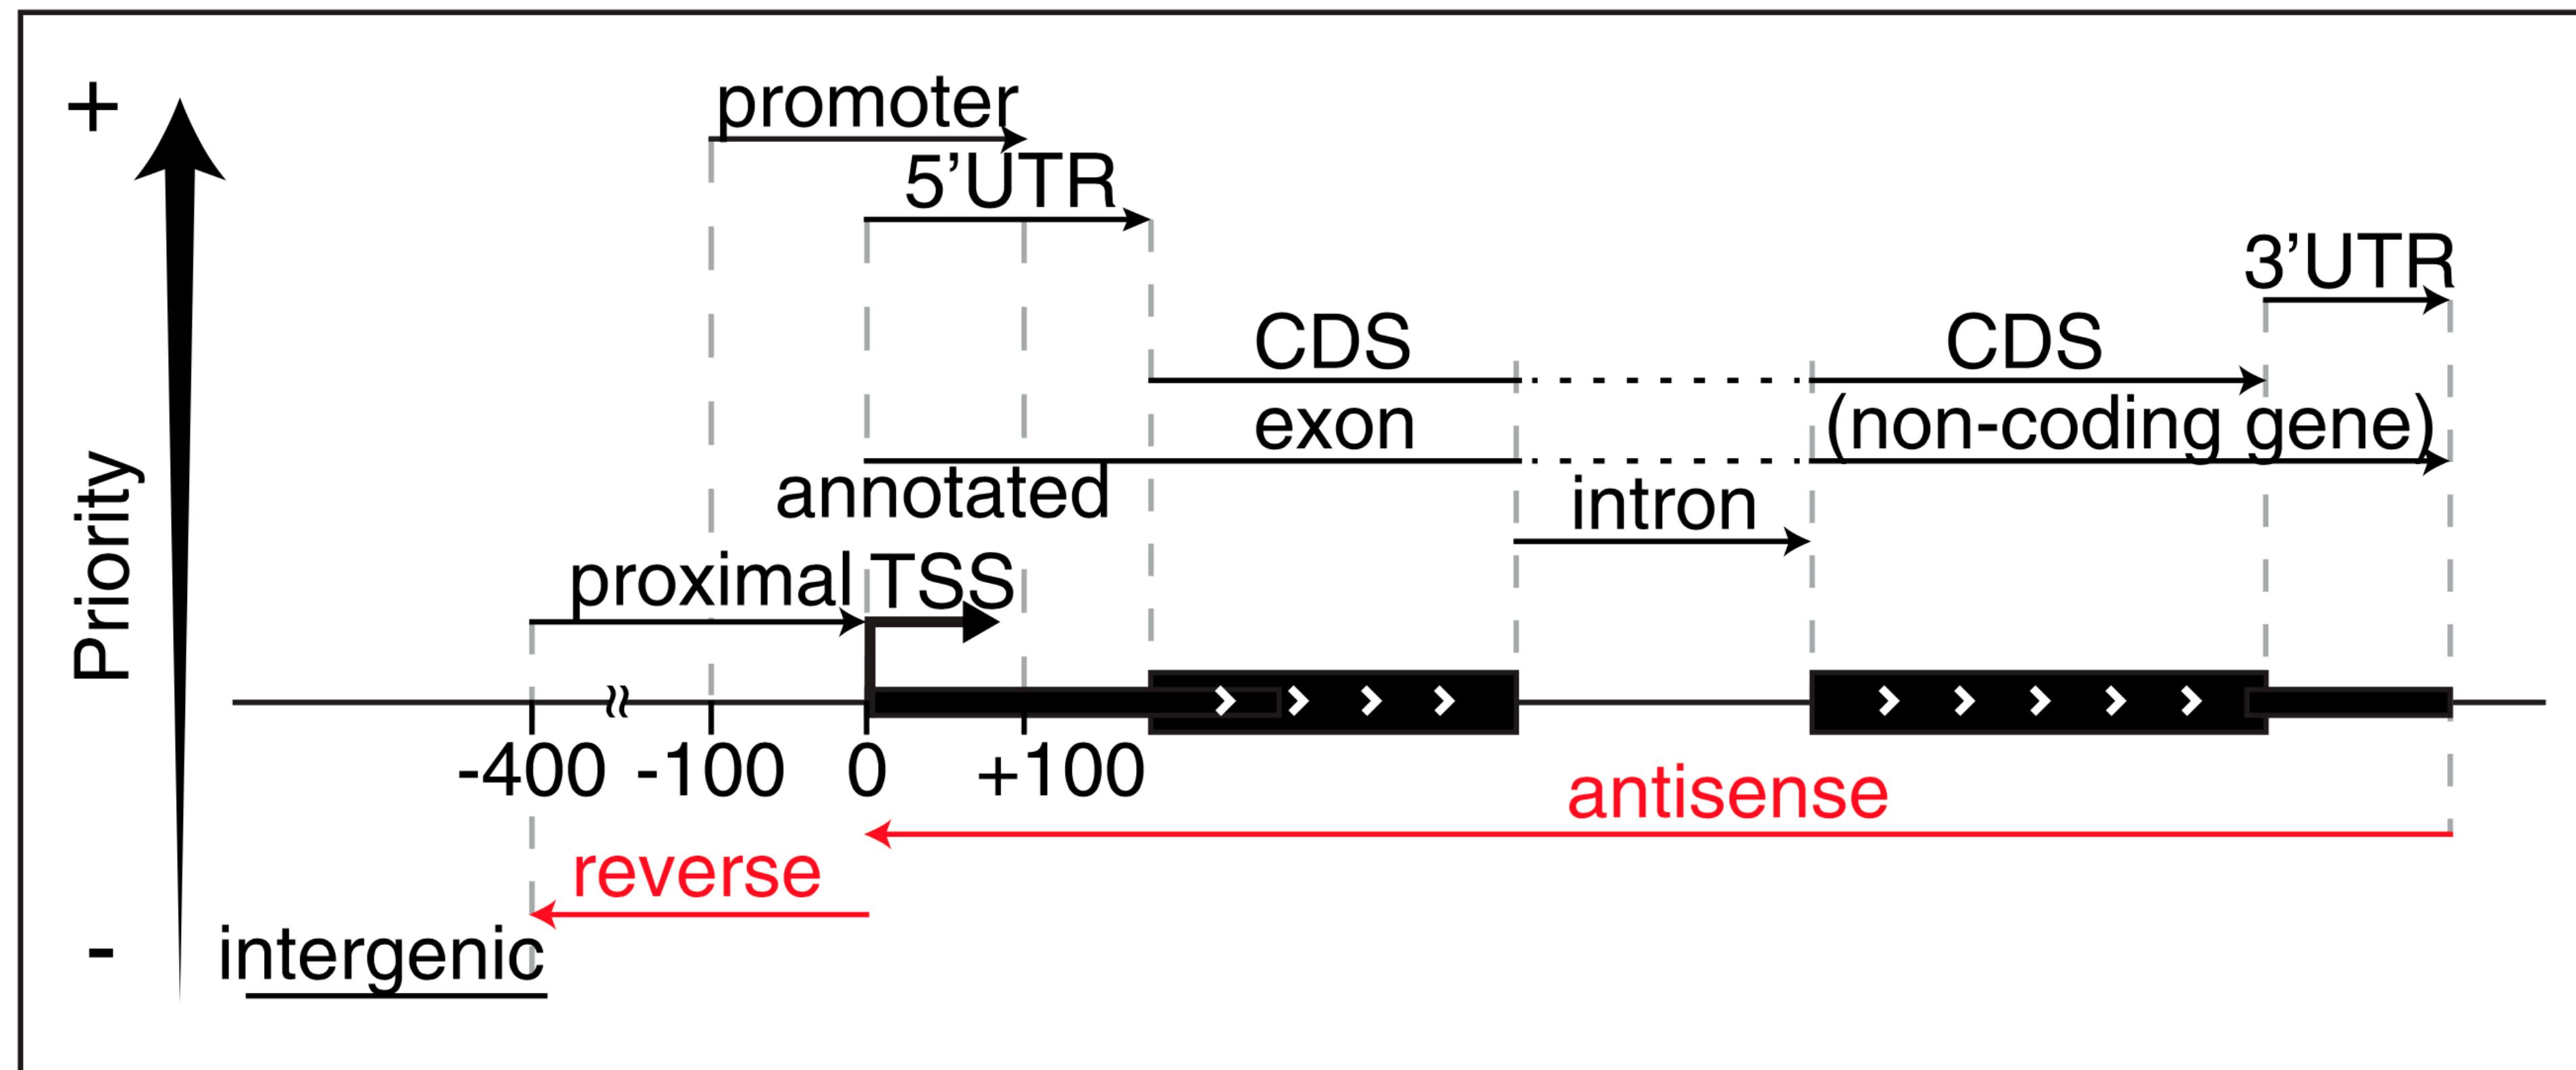
What CAGE looks like



What CAGE (really) looks like



CAGE annotation



Read the data!

=> <https://github.com/athieffry/HEH ATH 2021>

- Download the file **CAGE_TC_counts.xlsx**, the matrix of CAGE raw counts
- Read the file in R using `read_xlsx()` from the package `readxl`
 - call it “**TCs**”
 - this might require to install `readxl`
 - good idea: have a quick look at the file first
- What data structure resulted from the `read_xlsx()` import?
- How many samples are there?
- How many CAGE TCs are there?

Sanity checks

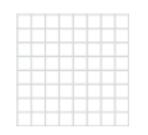
- It's always a good idea to sanity check data you receive, whoever they come from!
- How many tags (reads) were mapped in each sample?
 - use `select()` to remove the `TC_id` and `colSums()`
 - use the `magrittr` pipe
- Do we have TCs with non attributed (`NA`) expression?
 - tip: `is.na()` tests whether a value is `NA`
- Try `sum(c(1:3, NA))`
 - What happened?
 - What do you deduce with regards to the 2 previous exercices?
 - Can you fix it? Tip: look at the Vignette of `sum()`
- Get an overview of the TCs expression distribution
 - Use `summary()`

ggplot2

Overview of ggplot2 geometries

GRAPHICAL PRIMITIVES

```
a <- ggplot(economics, aes(date, unemploy))
b <- ggplot(seals, aes(x = long, y = lat))
```



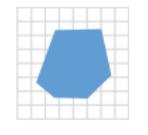
a + geom_blank() and **a + expand_limits()**
Ensure limits include values across all plots.



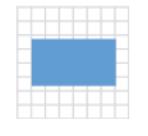
b + geom_curve(aes(yend = lat + 1, xend = long + 1, curvature = 1))
x, y, alpha, angle, color, curvature, linetype, size



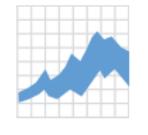
a + geom_path(lineend = "butt", linejoin = "round", linemitre = 1)
x, y, alpha, color, group, linetype, size



a + geom_polygon(aes(alpha = 50)) - x, y, alpha, color, fill, group, subgroup, linetype, size



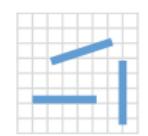
b + geom_rect(aes(xmin = long, ymin = lat, xmax = long + 1, ymax = lat + 1)) - xmax, xmin, ymax, ymin, alpha, color, fill, linetype, size



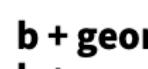
a + geom_ribbon(aes(ymax = unemploy - 900, ymin = unemploy + 900)) - x, ymax, ymin, alpha, color, fill, group, linetype, size

LINE SEGMENTS

common aesthetics: x, y, alpha, color, linetype, size



b + geom_abline(aes(intercept = 0, slope = 1))
b + geom_hline(aes(yintercept = lat))
b + geom_vline(aes(xintercept = long))



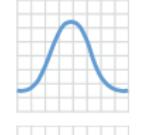
b + geom_segment(aes(yend = lat + 1, xend = long + 1))
b + geom_spoke(aes(angle = 1:1155, radius = 1))

ONE VARIABLE continuous

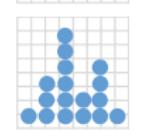
```
c <- ggplot(mpg, aes(hwy)); c2 <- ggplot(mpg)
```



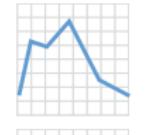
c + geom_area(stat = "bin")
x, y, alpha, color, fill, linetype, size



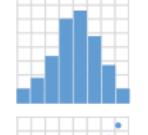
c + geom_density(kernel = "gaussian")
x, y, alpha, color, fill, group, linetype, size, weight



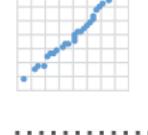
c + geom_dotplot()
x, y, alpha, color, fill



c + geom_freqpoly()
x, y, alpha, color, group, linetype, size



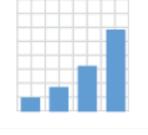
c + geom_histogram(binwidth = 5)
x, y, alpha, color, fill, linetype, size, weight



c2 + geom_qq(aes(sample = hwy))
x, y, alpha, color, fill, linetype, size, weight

discrete

```
d <- ggplot(mpg, aes(fl))
```



d + geom_bar()
x, alpha, color, fill, linetype, size, weight

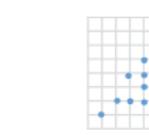
TWO VARIABLES

both continuous

```
e <- ggplot(mpg, aes(cty, hwy))
```



e + geom_label(aes(label = cty), nudge_x = 1, nudge_y = 1) - x, y, label, alpha, angle, color, family, fontface, hjust, lineheight, size, vjust



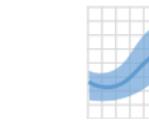
e + geom_point()
x, y, alpha, color, fill, shape, size, stroke



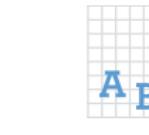
e + geom_quantile()
x, y, alpha, color, group, linetype, size, weight



e + geom_rug(sides = "bl")
x, y, alpha, color, linetype, size



e + geom_smooth(method = lm)
x, y, alpha, color, fill, group, linetype, size, weight



e + geom_text(aes(label = cty), nudge_x = 1, nudge_y = 1) - x, y, label, alpha, angle, color, family, fontface, hjust, lineheight, size, vjust

continuous bivariate distribution

```
h <- ggplot(diamonds, aes(carat, price))
```



h + geom_bin2d(binwidth = c(0.25, 500))
x, y, alpha, color, fill, linetype, size, weight



h + geom_density_2d()
x, y, alpha, color, group, linetype, size



h + geom_hex()
x, y, alpha, color, fill, size

continuous function

```
i <- ggplot(economics, aes(date, unemploy))
```



i + geom_area()
x, y, alpha, color, fill, linetype, size



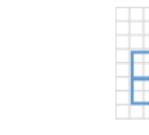
i + geom_line()
x, y, alpha, color, group, linetype, size



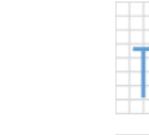
i + geom_step(direction = "hv")
x, y, alpha, color, group, linetype, size

visualizing error

```
df <- data.frame(grp = c("A", "B"), fit = 4:5, se = 1:2)
j <- ggplot(df, aes(grp, fit, ymin = fit - se, ymax = fit + se))
```



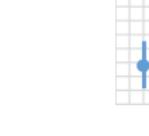
j + geom_crossbar(fatten = 2) - x, y, ymax, ymin, alpha, color, fill, group, linetype, size



j + geom_errorbar() - x, y, ymax, ymin, alpha, color, group, linetype, size
Also **geom_errorbarh()**.



j + geom_linerange()
x, ymin, ymax, alpha, color, group, linetype, size



j + geom_pointrange() - x, y, ymin, ymax, alpha, color, fill, group, linetype, shape, size

maps

```
data <- data.frame(murder = USArrests$Murder,
```

```
state = tolower(rownames(USArrests)))
```

```
map <- map_data("state")
```

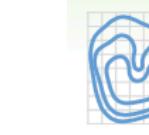
```
k <- ggplot(data, aes(fill = murder))
```



k + geom_map(aes(map_id = state), map = map) + expand_limits(x = map\$long, y = map\$lat)
map_id, alpha, color, fill, linetype, size

THREE VARIABLES

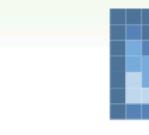
```
seals$z <- with(seals, sqrt(delta_long^2 + delta_lat^2)); l <- ggplot(seals, aes(long, lat))
```



l + geom_contour(aes(z = z))
x, y, z, alpha, color, group, linetype, size, weight



l + geom_contour_filled(aes(fill = z))
x, y, alpha, color, fill, group, linetype, size, subgroup



l + geom_raster(aes(fill = z), hjust = 0.5, vjust = 0.5, interpolate = FALSE)
x, y, alpha, fill



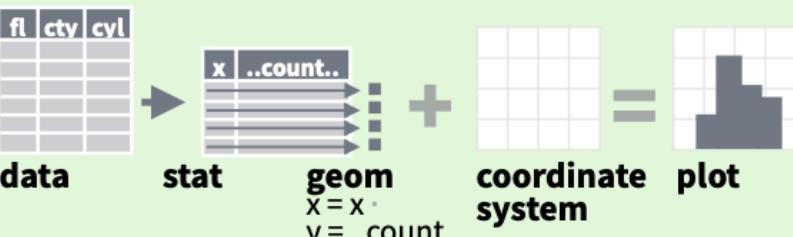
l + geom_tile(aes(fill = z))
x, y, alpha, color, fill, linetype, size, width

Overview of ggplot2 geometries

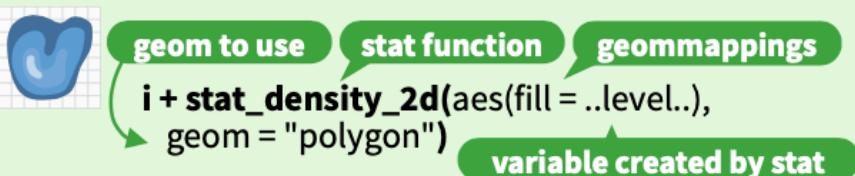
Stats

An alternative way to build a layer.

A stat builds new variables to plot (e.g., count, prop).



Visualize a stat by changing the default stat of a geom function, `geom_bar(stat="count")` or by using a stat function, `stat_count(geom="bar")`, which calls a default geom to make a layer (equivalent to a geom function). Use `..name..` syntax to map stat variables to aesthetics.



```

c + stat_bin(binwidth = 1, boundary = 10)
x, y | ..count.., ..density.., ..ndensity..
c + stat_count(width = 1) x, y | ..count.., ..prop..
c + stat_density(adjust = 1, kernel = "gaussian")
x, y | ..count.., ..density.., ..scaled..
  
```

```

e + stat_bin_2d(bins = 30, drop = T)
x, y, fill | ..count.., ..density..
e + stat_bin_hex(bins = 30) x, y, fill | ..count.., ..density..
  
```

```

e + stat_density_2d(contour = TRUE, n = 100)
x, y, color, size | ..level..
e + stat_ellipse(level = 0.95, segments = 51, type = "t")
  
```

```

l + stat_contour(aes(z = z)) x, y, z, order | ..level..
l + stat_summary_hex(aes(z = z), bins = 30, fun = max)
x, y, z, fill | ..value..
  
```

```

l + stat_summary_2d(aes(z = z), bins = 30, fun = mean)
x, y, z, fill | ..value..
f + stat_boxplot(coef = 1.5)
x, y | ..lower.., ..middle.., ..upper.., ..width.., ..ymin.., ..ymax..
  
```

```

f + stat_ydensity(kernel = "gaussian", scale = "area") x, y
| ..density.., ..scaled.., ..count.., ..n.., ..violinwidth.., ..width..
e + stat_ecdf(n = 40) x, y | ..x.., ..y..
  
```

```

e + stat_quantile(quantiles = c(0.1, 0.9),
formula = y ~ log(x), method = "rq") x, y | ..quantile..
e + stat_smooth(method = "lm", formula = y ~ x, se = T,
level = 0.95) x, y | ..se.., ..x.., ..y.., ..ymin.., ..ymax..
  
```

```

ggplot() + xlim(-5, 5) + stat_function(fun = dnorm,
n = 20, geom = "point") x | ..x.., ..y..
  
```

```

ggplot() + stat_qq(aes(sample = 1:100))
x, y, sample | ..sample.., ..theoretical..
  
```

```

e + stat_sum() x, y, size | ..n.., ..prop..
e + stat_summary(fun.data = "mean_cl_boot")
  
```

```

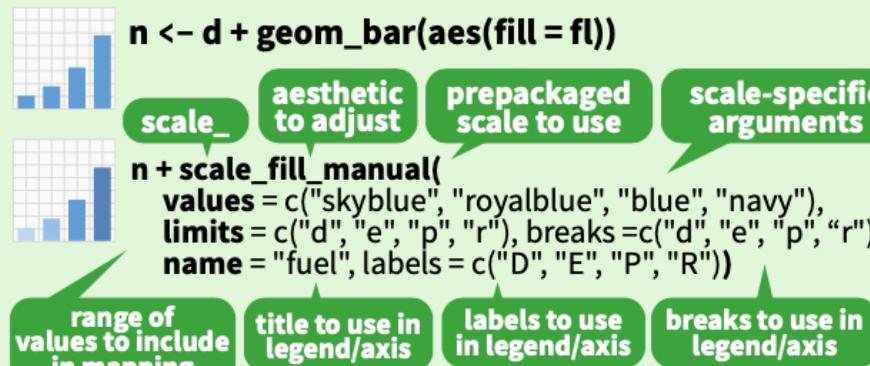
h + stat_summary_bin(fun = "mean", geom = "bar")
e + stat_identity()
e + stat_unique()
  
```



Scales

Override defaults with `scales` package.

Scales map data values to the visual values of an aesthetic. To change a mapping, add a new scale.



GENERAL PURPOSE SCALES

Use with most aesthetics

```

scale_*_continuous() - Map cont' values to visual ones.
scale_*_discrete() - Map discrete values to visual ones.
scale_*_binned() - Map continuous values to discrete bins.
scale_*_identity() - Use data values as visual ones.
scale_*_manual(values = c()) - Map discrete values to manually chosen visual ones.
scale_*_date(date_labels = "%m/%d"),
date_breaks = "2 weeks") - Treat data values as dates.
scale_*_datetime() - Treat data values as date times.
Same as scale_*_date(). See ?strptime for label formats.
  
```

X & Y LOCATION SCALES

Use with x or y aesthetics (x shown here)

```

scale_x_log10() - Plot x on log10 scale.
scale_x_reverse() - Reverse the direction of the x axis.
scale_x_sqrt() - Plot x on square root scale.
  
```

COLOR AND FILL SCALES (DISCRETE)

```

n + scale_fill_brewer(palette = "Blues")
For palette choices:
RColorBrewer::display.brewer.all()
n + scale_fill_grey(start = 0.2,
end = 0.8, na.value = "red")
  
```

COLOR AND FILL SCALES (CONTINUOUS)

```

o <- c + geom_dotplot(aes(fill = ..x..))
o + scale_fill_distiller(palette = "Blues")
o + scale_fill_gradient(low = "red", high = "yellow")
o + scale_fill_gradient2(low = "red", high = "blue",
mid = "white", midpoint = 25)
o + scale_fill_gradientn(colors = topo.colors(6))
Also: rainbow(), heat.colors(), terrain.colors(),
cm.colors(), RColorBrewer::brewer.pal()
  
```

SHAPE AND SIZE SCALES

```

p <- e + geom_point(aes(shape = fl, size = cyl))
  
```

```

p + scale_shape() + scale_size()
p + scale_shape_manual(values = c(3:7))
p + scale_radius(range = c(1,6))
p + scale_size_area(max_size = 6)
  
```

Coordinate Systems

`r <- d + geom_bar()`

`r + coord_cartesian(xlim = c(0, 5))` - xlim, ylim
The default cartesian coordinate system.

`r + coord_fixed(ratio = 1/2)`
ratio, xlim, ylim - Cartesian coordinates with fixed aspect ratio between x and y units.

`ggplot(mpg, aes(y = fl)) + geom_bar()`
Flip cartesian coordinates by switching x and y aesthetic mappings.

`r + coord_polar(theta = "x", direction = 1)`
theta, start, direction - Polar coordinates.

`r + coord_trans(y = "sqrt")` - x, y, xlim, ylim
Transformed cartesian coordinates. Set xtrans and ytrans to the name of a window function.

`π + coord_quickmap()`
`π + coord_map(projection = "ortho", orientation = c(41, -74, 0))` - projection, xlim, ylim
Map projections from the mapproj package (mercator (default), azequalarea, lagrange, etc.).

Position Adjustments

Position adjustments determine how to arrange geoms that would otherwise occupy the same space.

```
s <- ggplot(mpg, aes(fl, fill = drv))
```

`s + geom_bar(position = "dodge")`
Arrange elements side by side.

`s + geom_bar(position = "fill")`
Stack elements on top of one another, normalize height.

`e + geom_point(position = "jitter")`
Add random noise to X and Y position of each element to avoid overplotting.

`e + geom_label(position = "nudge")`
Nudge labels away from points.

`s + geom_bar(position = "stack")`
Stack elements on top of one another.

Each position adjustment can be recast as a function with manual `width` and `height` arguments:

```
s + geom_bar(position = position_dodge(width = 1))
```

Themes

`r + theme_bw()`
White background with grid lines.

`r + theme_gray()`
Grey background (default theme).

`r + theme_dark()`
Dark for contrast.

`r + theme_classic()`

`r + theme_light()`

`r + theme_linedraw()`

`r + theme_minimal()`
Minimal theme.

`r + theme_void()`
Empty theme.

`r + theme()` Customize aspects of the theme such as axis, legend, panel, and facet properties.

```
r + ggtitle("Title") + theme(plot.title.position = "plot")
r + theme(panel.background = element_rect(fill = "blue"))
```

Faceting

Facets divide a plot into subplots based on the values of one or more discrete variables.

```
t <- ggplot(mpg, aes(cty, hwy)) + geom_point()
```

`t + facet_grid(cols = vars(fl))`
Facet into columns based on fl.

`t + facet_grid(rows = vars(year))`
Facet into rows based on year.

`t + facet_grid(rows = vars(year), cols = vars(fl))`
Facet into both rows and columns.

`t + facet_wrap(vars(fl))`
Wrap facets into a rectangular layout.

Set `scales` to let axis limits vary across facets.

`t + facet_grid(rows = vars(drv), cols = vars(fl), scales = "free")`
x and y axis limits adjust to individual facets:
"free_x" - x axis limits adjust
"free_y" - y axis limits adjust

Set `labeler` to adjust facet label:

`t + facet_grid(cols = vars(fl), labeler = label_both)`

`fl: c fl: d fl: e fl: p fl: r`

`t + facet_grid(rows = vars(fl), labeler = label_bquote(alpha ^ .(fl)))`

`αc αd αe αp αr`

Labels and Legends

Use `labs()` to label the elements of your plot.

`t + labs(x = "New x axis label", y = "New y axis label", title = "Add a title above the plot", subtitle = "Add a subtitle below title", caption = "Add a caption below plot", alt = "Add alt text to the plot", <AES> = "New <AES> legend title")`

`t + annotate(geom = "text", x = 8, y = 9, label = "A")`
Places a geom with manually selected aesthetics.

`p + guides(x = guide_axis(n.dodge = 2))` Avoid crowded or overlapping labels with `guide_axis(n.dodge` or `angle`).

`n + guides(fill = "none")` Set legend type for each aesthetic: colorbar, legend, or none (no legend).

`n + theme(legend.position = "bottom")` Place legend at "bottom", "top", "left", or "right".

`n + scale_fill_discrete(name = "Title", labels = c("A", "B", "C", "D", "E"))`
Set legend title and labels with a scale function.

Zooming

Without clipping (preferred):

`t + coord_cartesian(xlim = c(0, 100), ylim = c(10, 20))`

With clipping (removes unseen data points):

`t + xlim(0, 100) + ylim(10, 20)`

`t + scale_x_continuous(limits = c(0, 100)) + scale_y_continuous(limits = c(0, 100))`



Our first plot!

- Looking at numbers in a table is a difficult mental exercise: a figure is needed!
- Reproduce the following code and try to dissect it:

```
26 # c. our first plot!
27 TCs %>%
28   select(-TC_id) %>%
29   colSums() %>%
30   enframe() %>%
31   ggplot(aes(x=name, y=value)) +
32     geom_col() +
33     labs(x='sample', y='library size')
```

Tips: run it line by line and see the intermediate steps.

Let's go through it.

- Here we use `select()` to remove (- sign) the column names `TC_id`

```
26 # c. our first plot!
27 TCs %>%
28   select(-TC_id)
```

- The only reason for the pipe here is to better understand what is the intput once we have a larger block of code

Let's go through it.

- Now that we only have numerical values left, we can compute the column-wise sums with `colSums()`

```
26 # c. our first plot!
27 TCs %>%
28   select(-TC_id) %>%
29   colSums()
```

Let's go through it.

- `ggplot2` only takes data frames as input, but the output of `colSums()` is a vector

```
26 # c. our first plot!
27 TCs %>%
28   select(-TC_id) %>%
29   colSums() %>%
30   enframe()
```

- The `enframe()` function nicely coerce the vector object into a data frame
- `as.data.frame()` could be used too, but the formatting is not as nice

Let's go through it.

- Now we pass the data frame to `ggplot()`
- `aes()` is where we define the plot-wide aesthetic and map variables to features of the plot. Here we map the X-axis to the sample names, and Y-axis to the value (those are default names of the `enframe()` function)
- Running `ggplot()` alone will just create a canvas, not geometric layer will be there

```
26 # c. our first plot!
27 TCs %>%
28   select(-TC_id) %>%
29   colSums() %>%
30   enframe() %>%
31   ggplot(aes(x=name, y=value))
```

Let's go through it.

- To add any layer to a `ggplot()` figure, we use ‘+’
- Here we use a column geometry by using `geom_col()`
- `geom_col()` will **inherit** the aesthetics that we declared in the main call of `ggplot()`

```
26 # c. our first plot!
27 TCs %>%
28   select(-TC_id) %>%
29   colSums() %>%
30   enframe() %>%
31   ggplot(aes(x=name, y=value)) +
32     |     | geom_col()
```

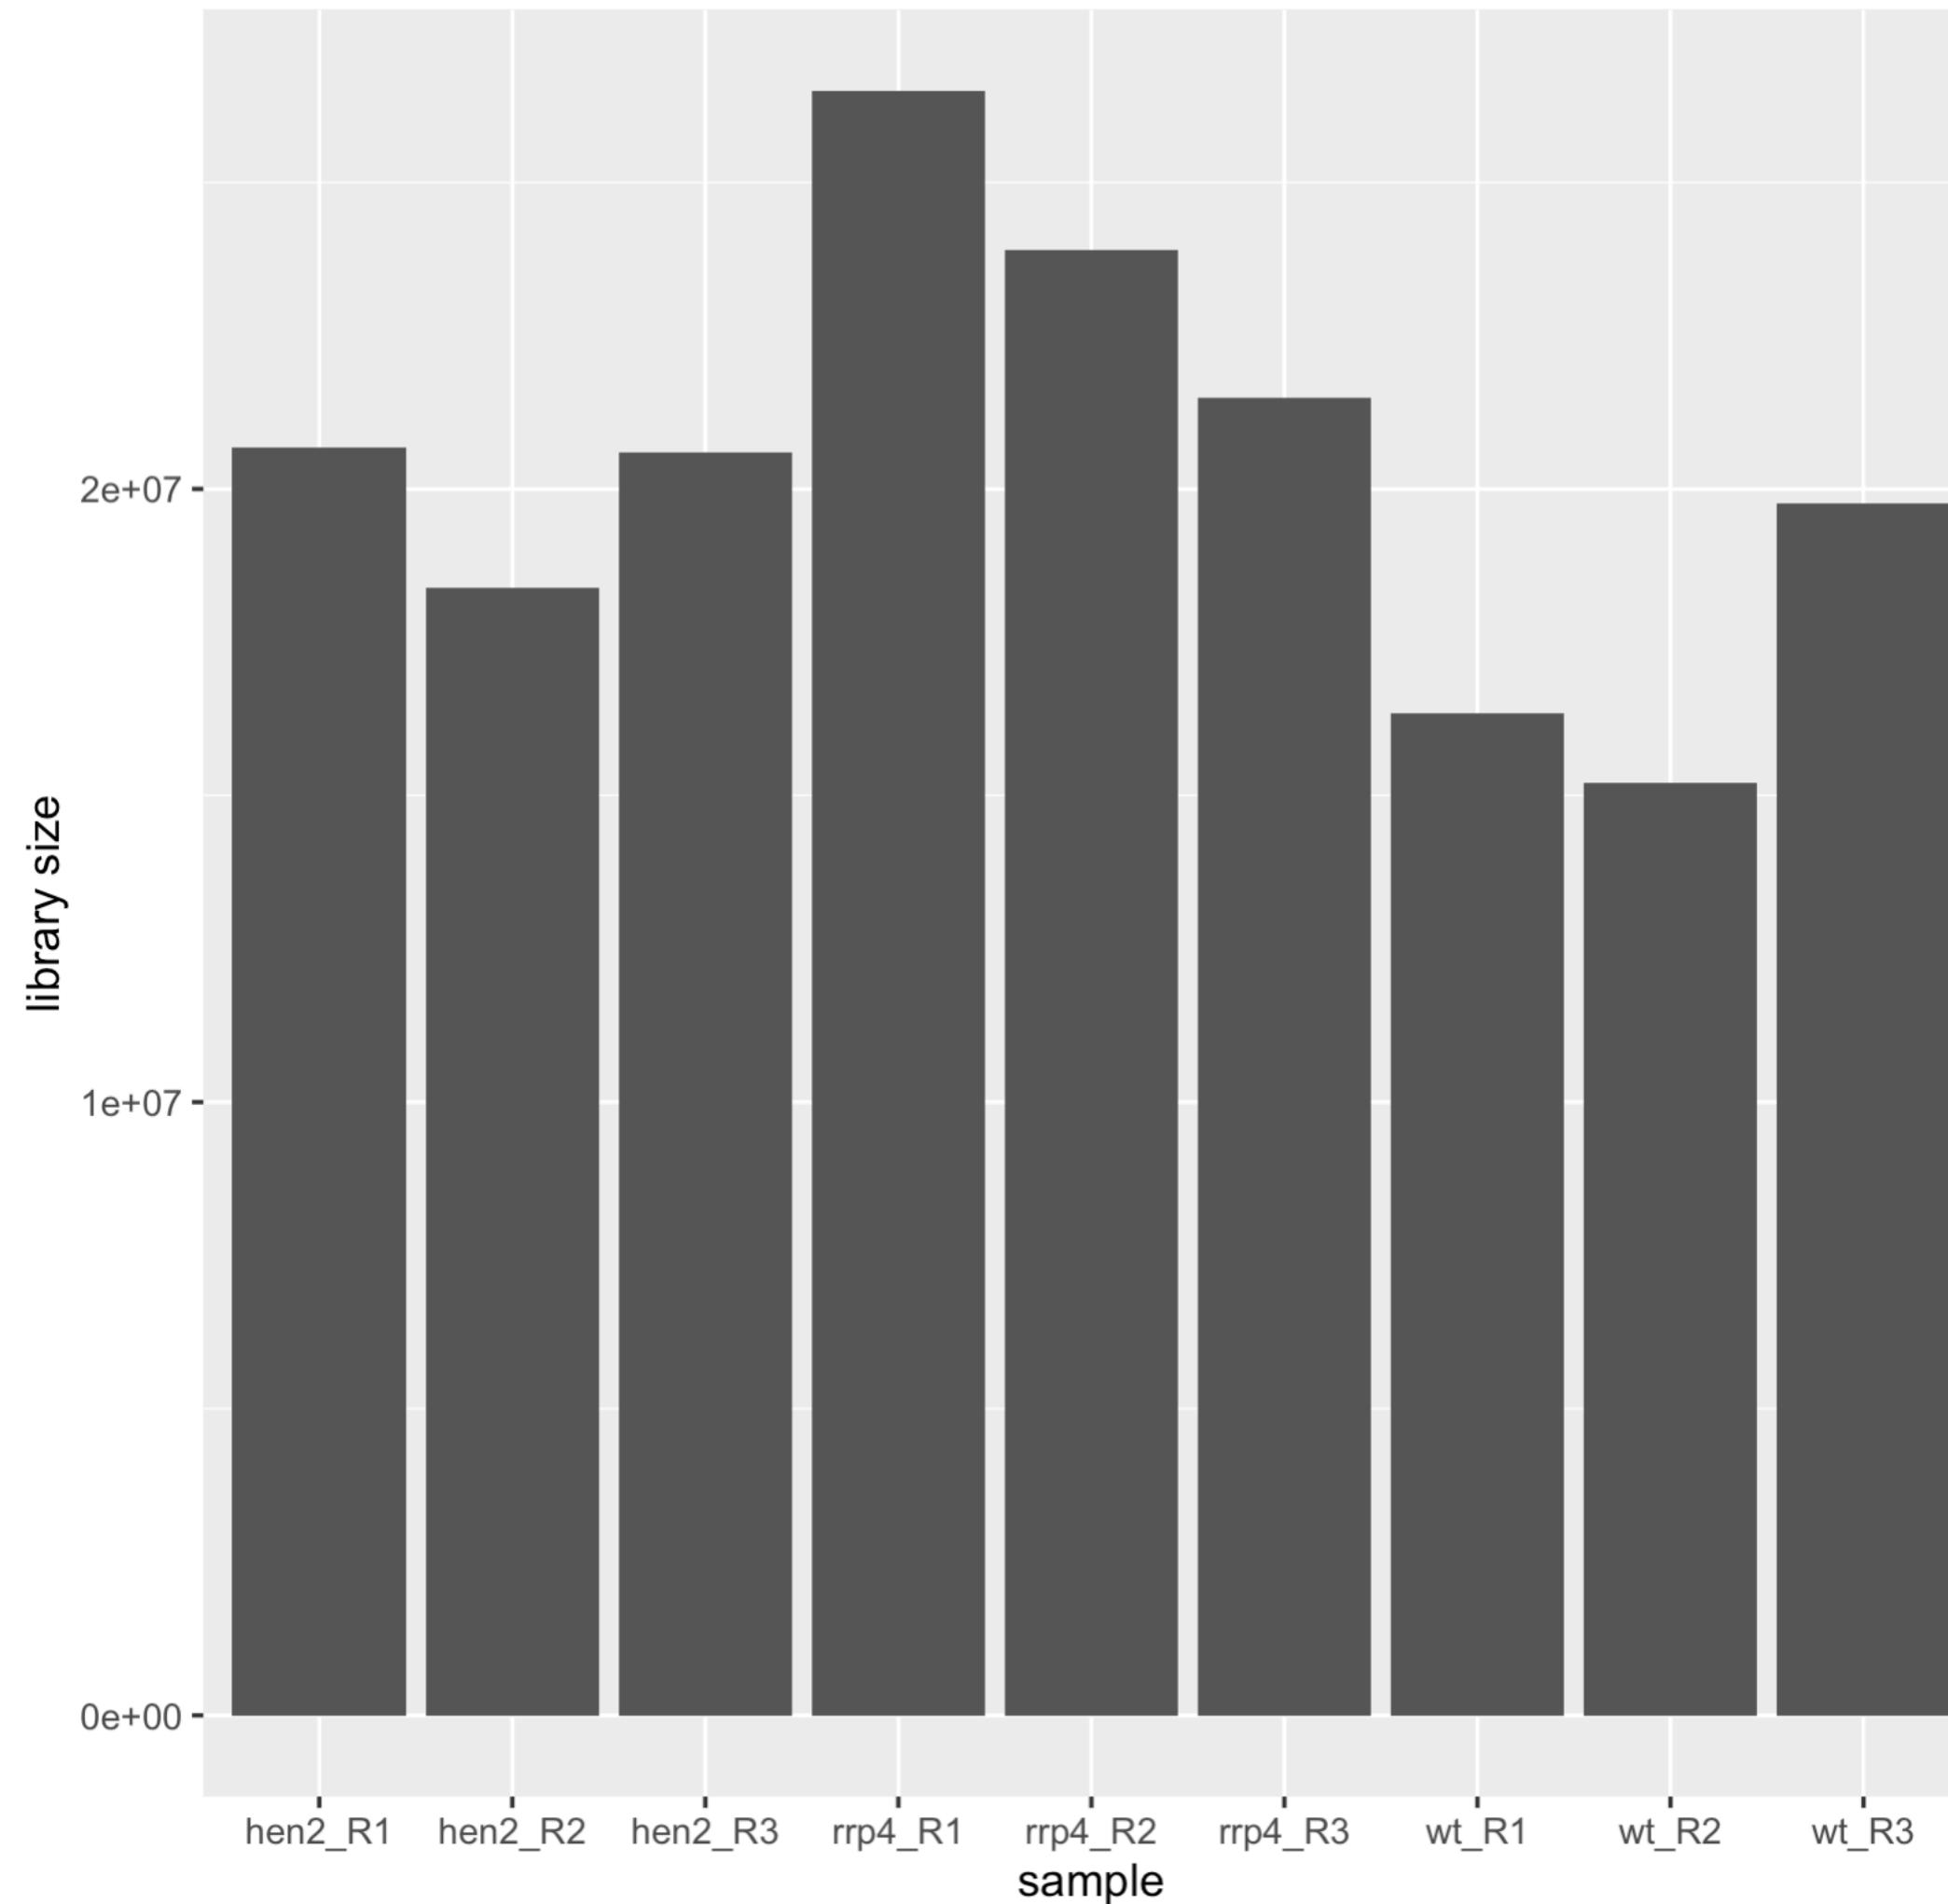
Let's go through it.

- Let's rename the axis labels to better fit the data
- We add another layer (+) called `labs()`
- Other parameters of `labs()` are title, subtitle, caption, x, y, ...

```
26 # c. our first plot!
27 TCs %>%
28   select(-TC_id) %>%
29   colSums() %>%
30   enframe() %>%
31   ggplot(aes(x=name, y=value)) +
32     geom_col() +
33     labs(x='sample', y='library size')
```

Let's go through it.

- Et voilà.

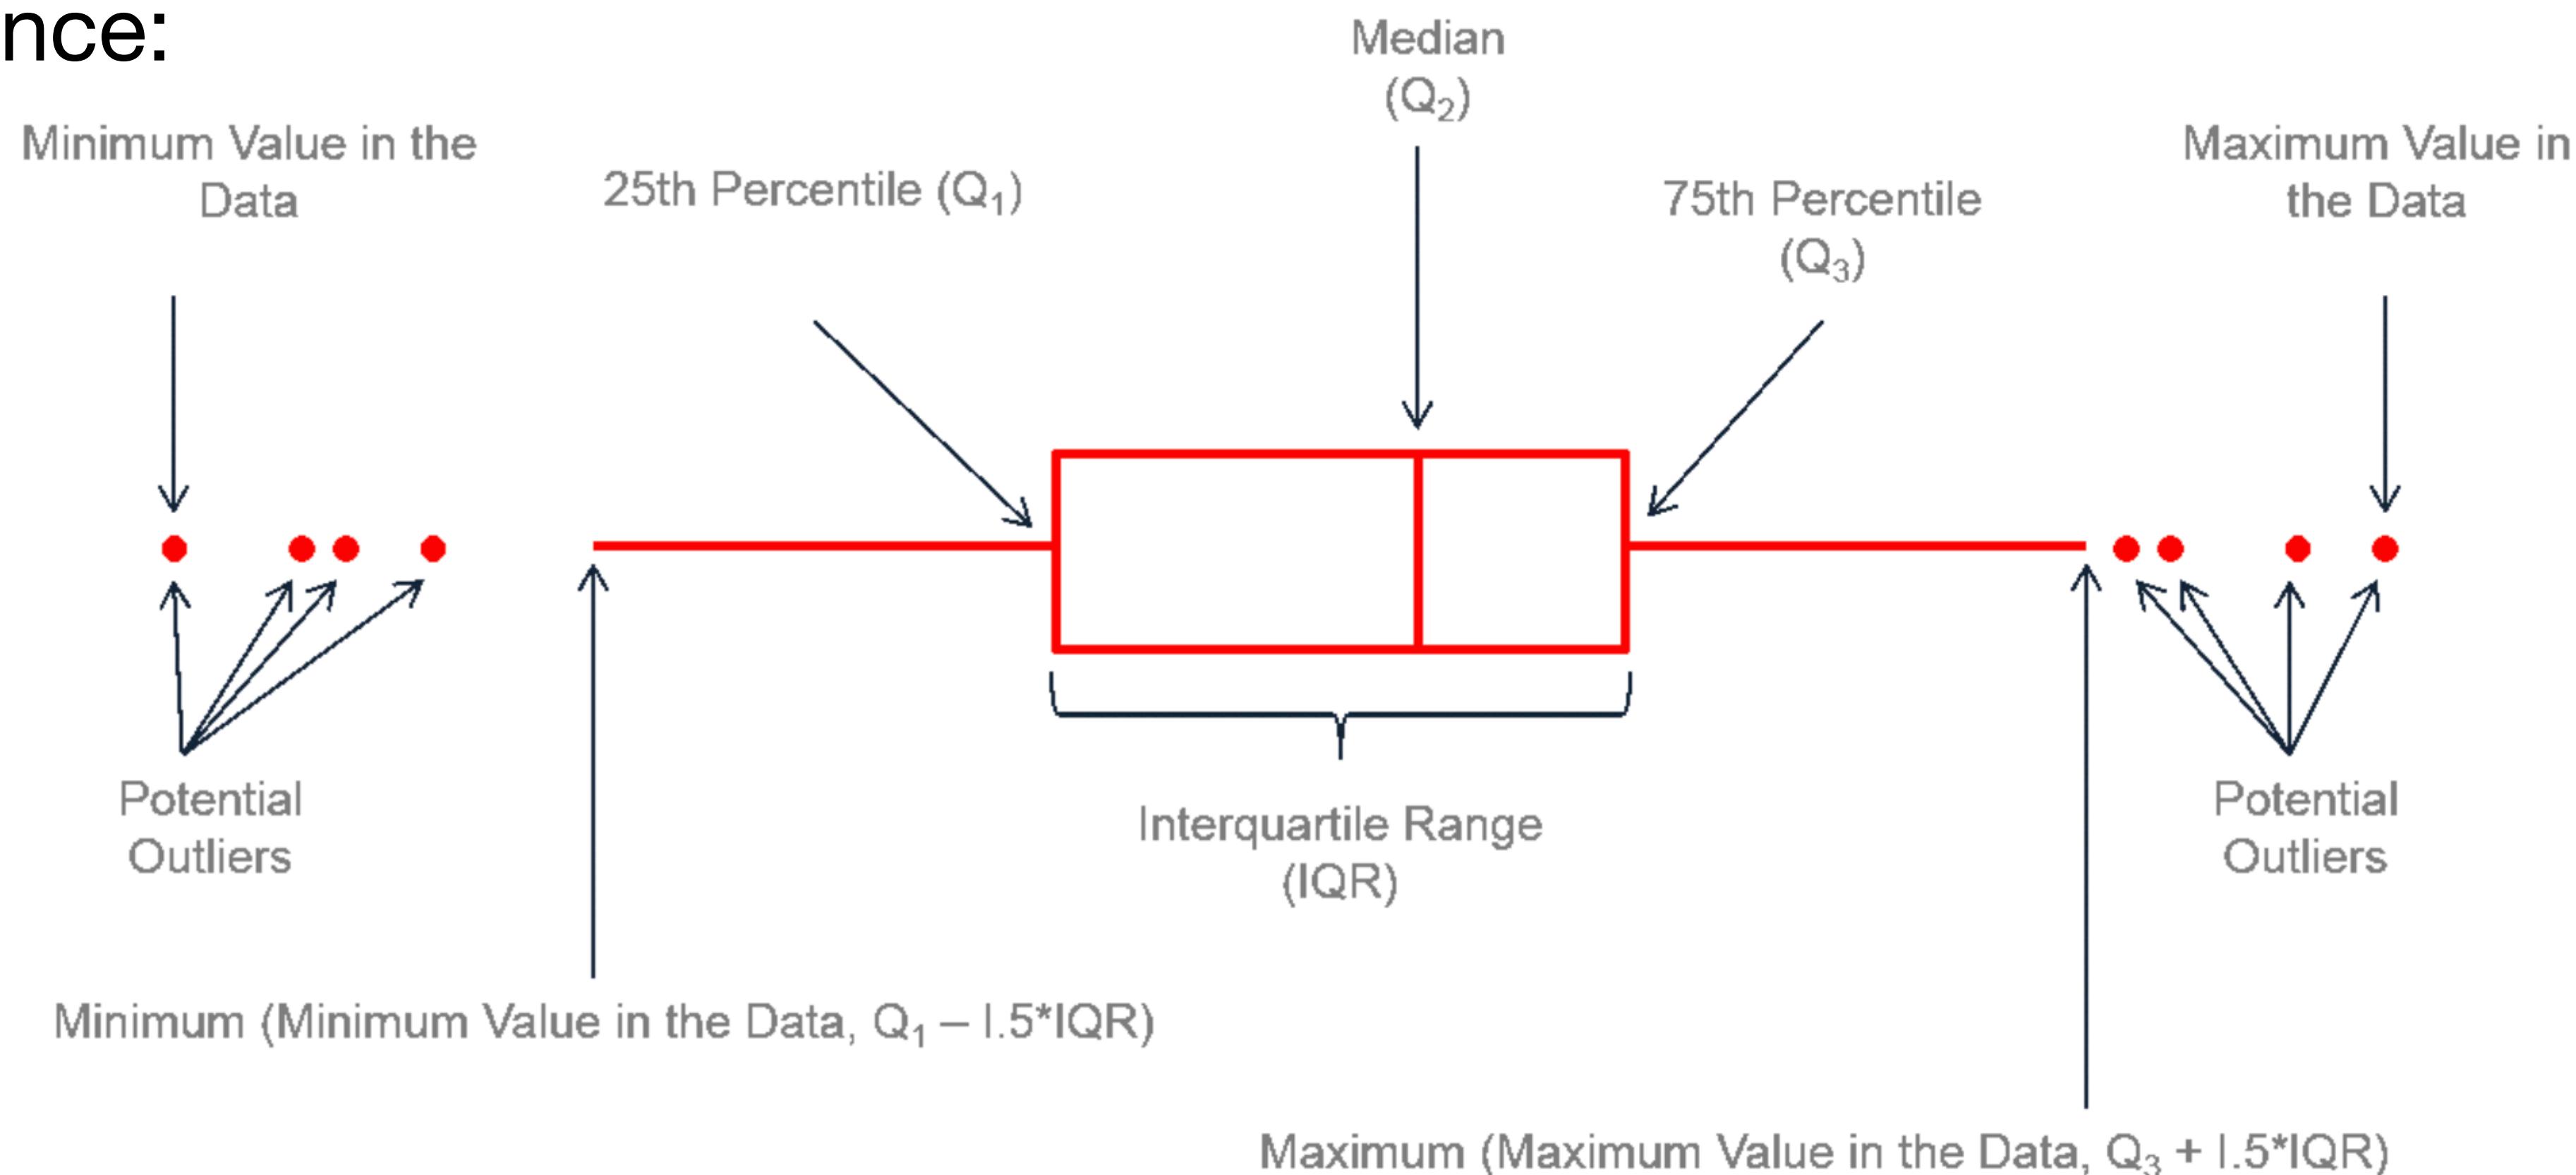


- `ggplot()` is a big library
- Many, many stuff are possible and customisable
- It relies on the **grammar of graphics**
- **MUST HAVE:**
ggplot2 cheat sheet

👉 www.rstudio.com/resources/cheatsheets/

Expression distribution: box plot

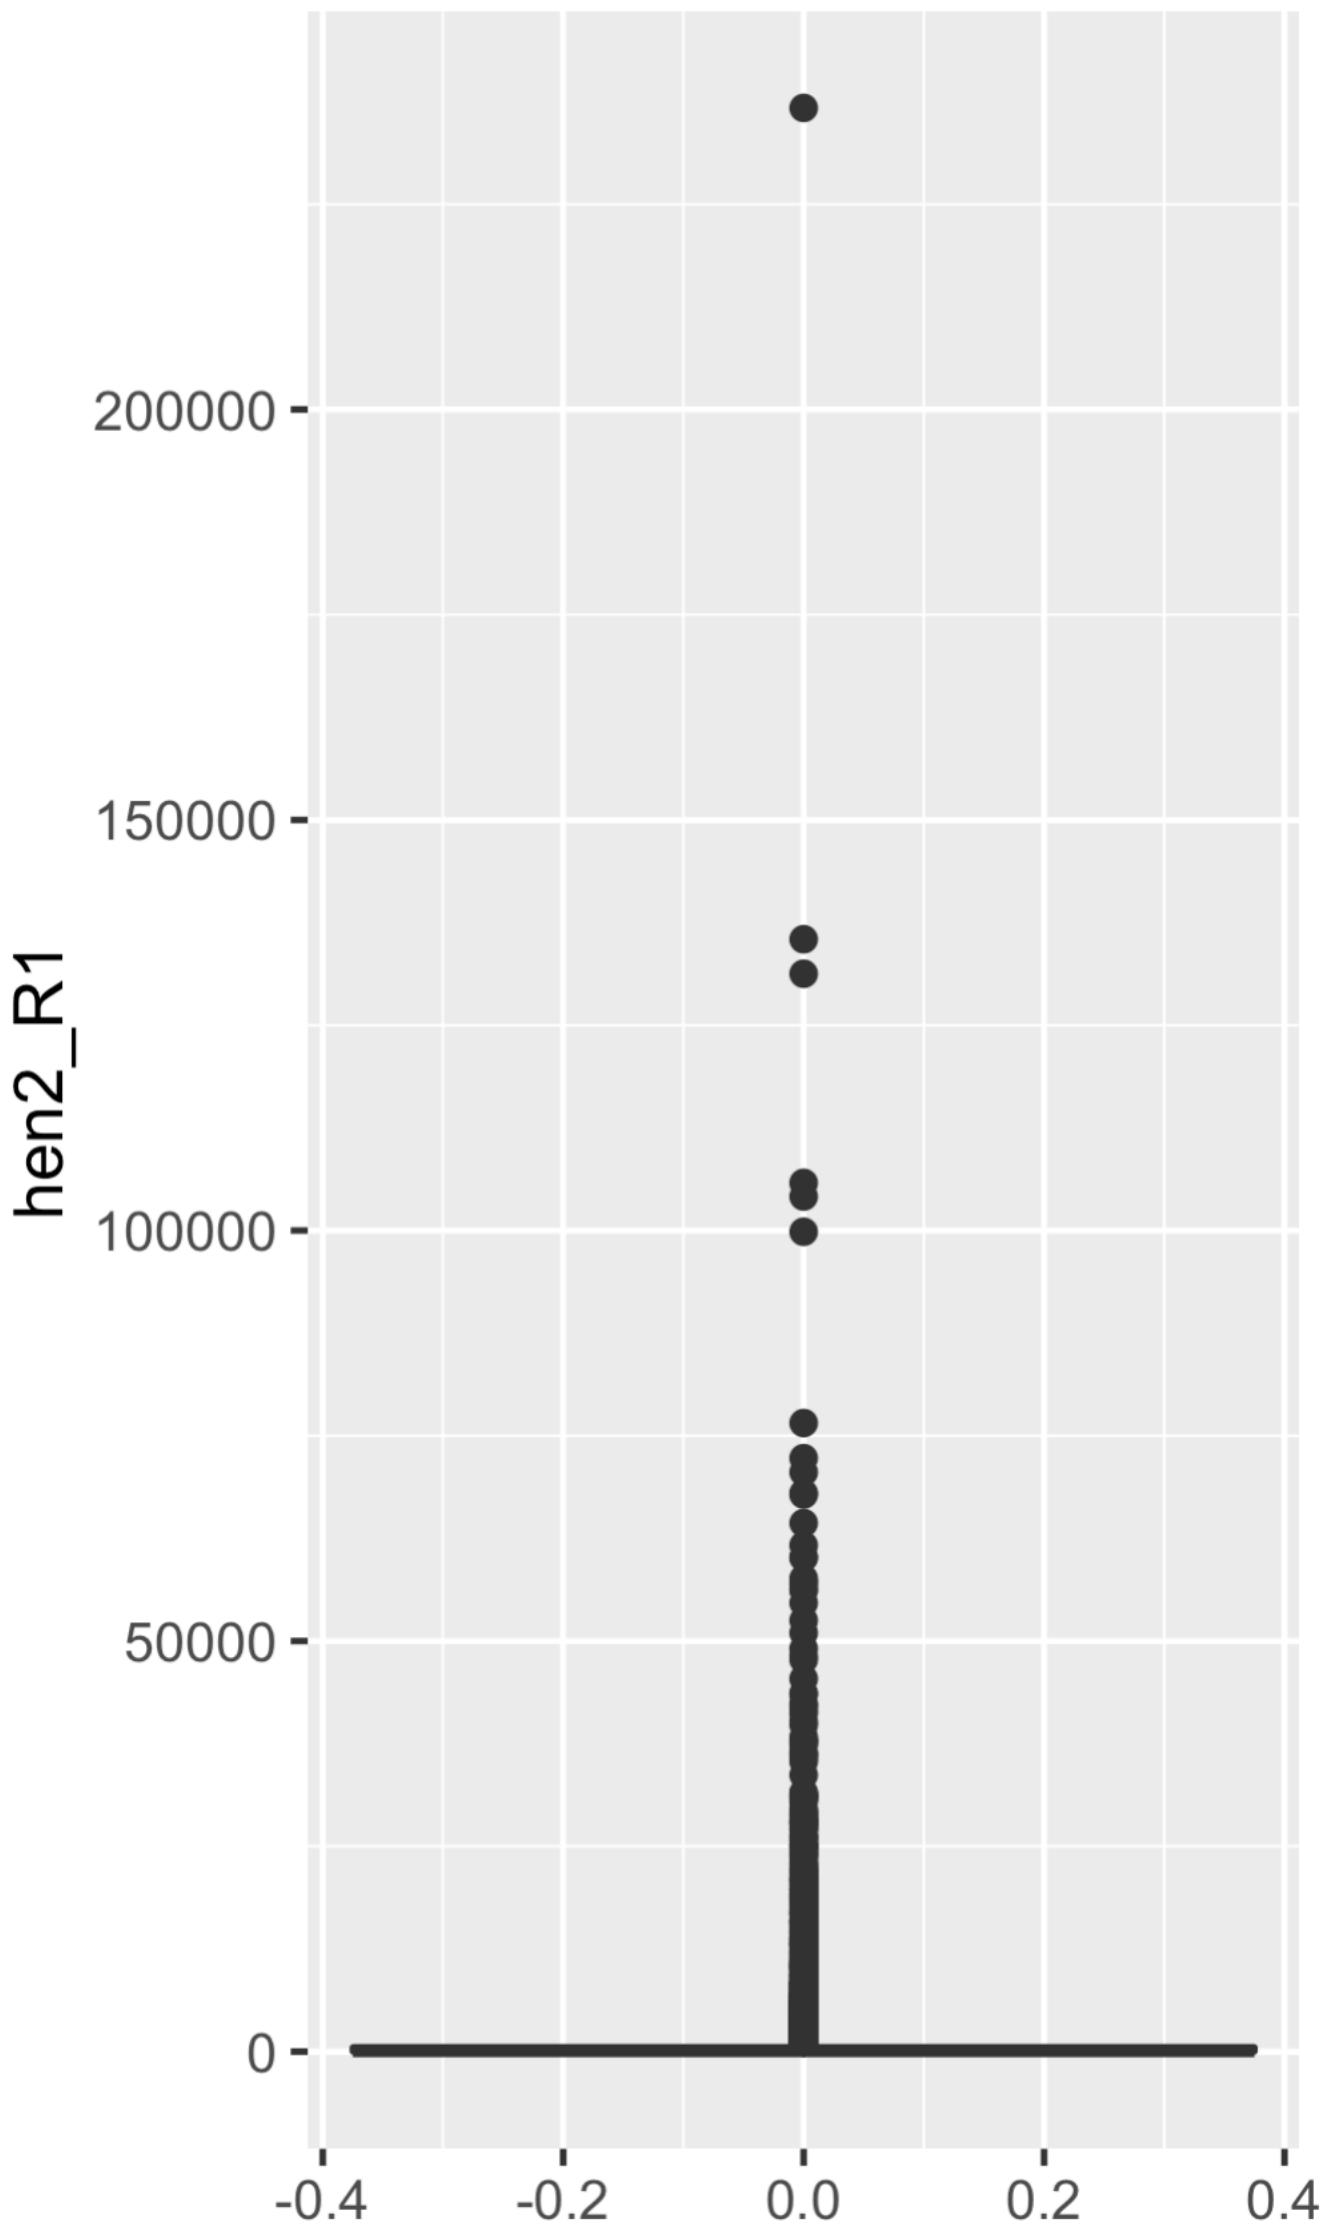
- A simple yet very informative plot for investigating distribution is the BOX PLOT
- Read more about boxplots <https://www.data-to-viz.com/caveat/boxplot.html>
- In essence:



Our first box plot

```
ggplot(data=TCs, aes(y=hen2_R1)) +  
  geom_boxplot()
```

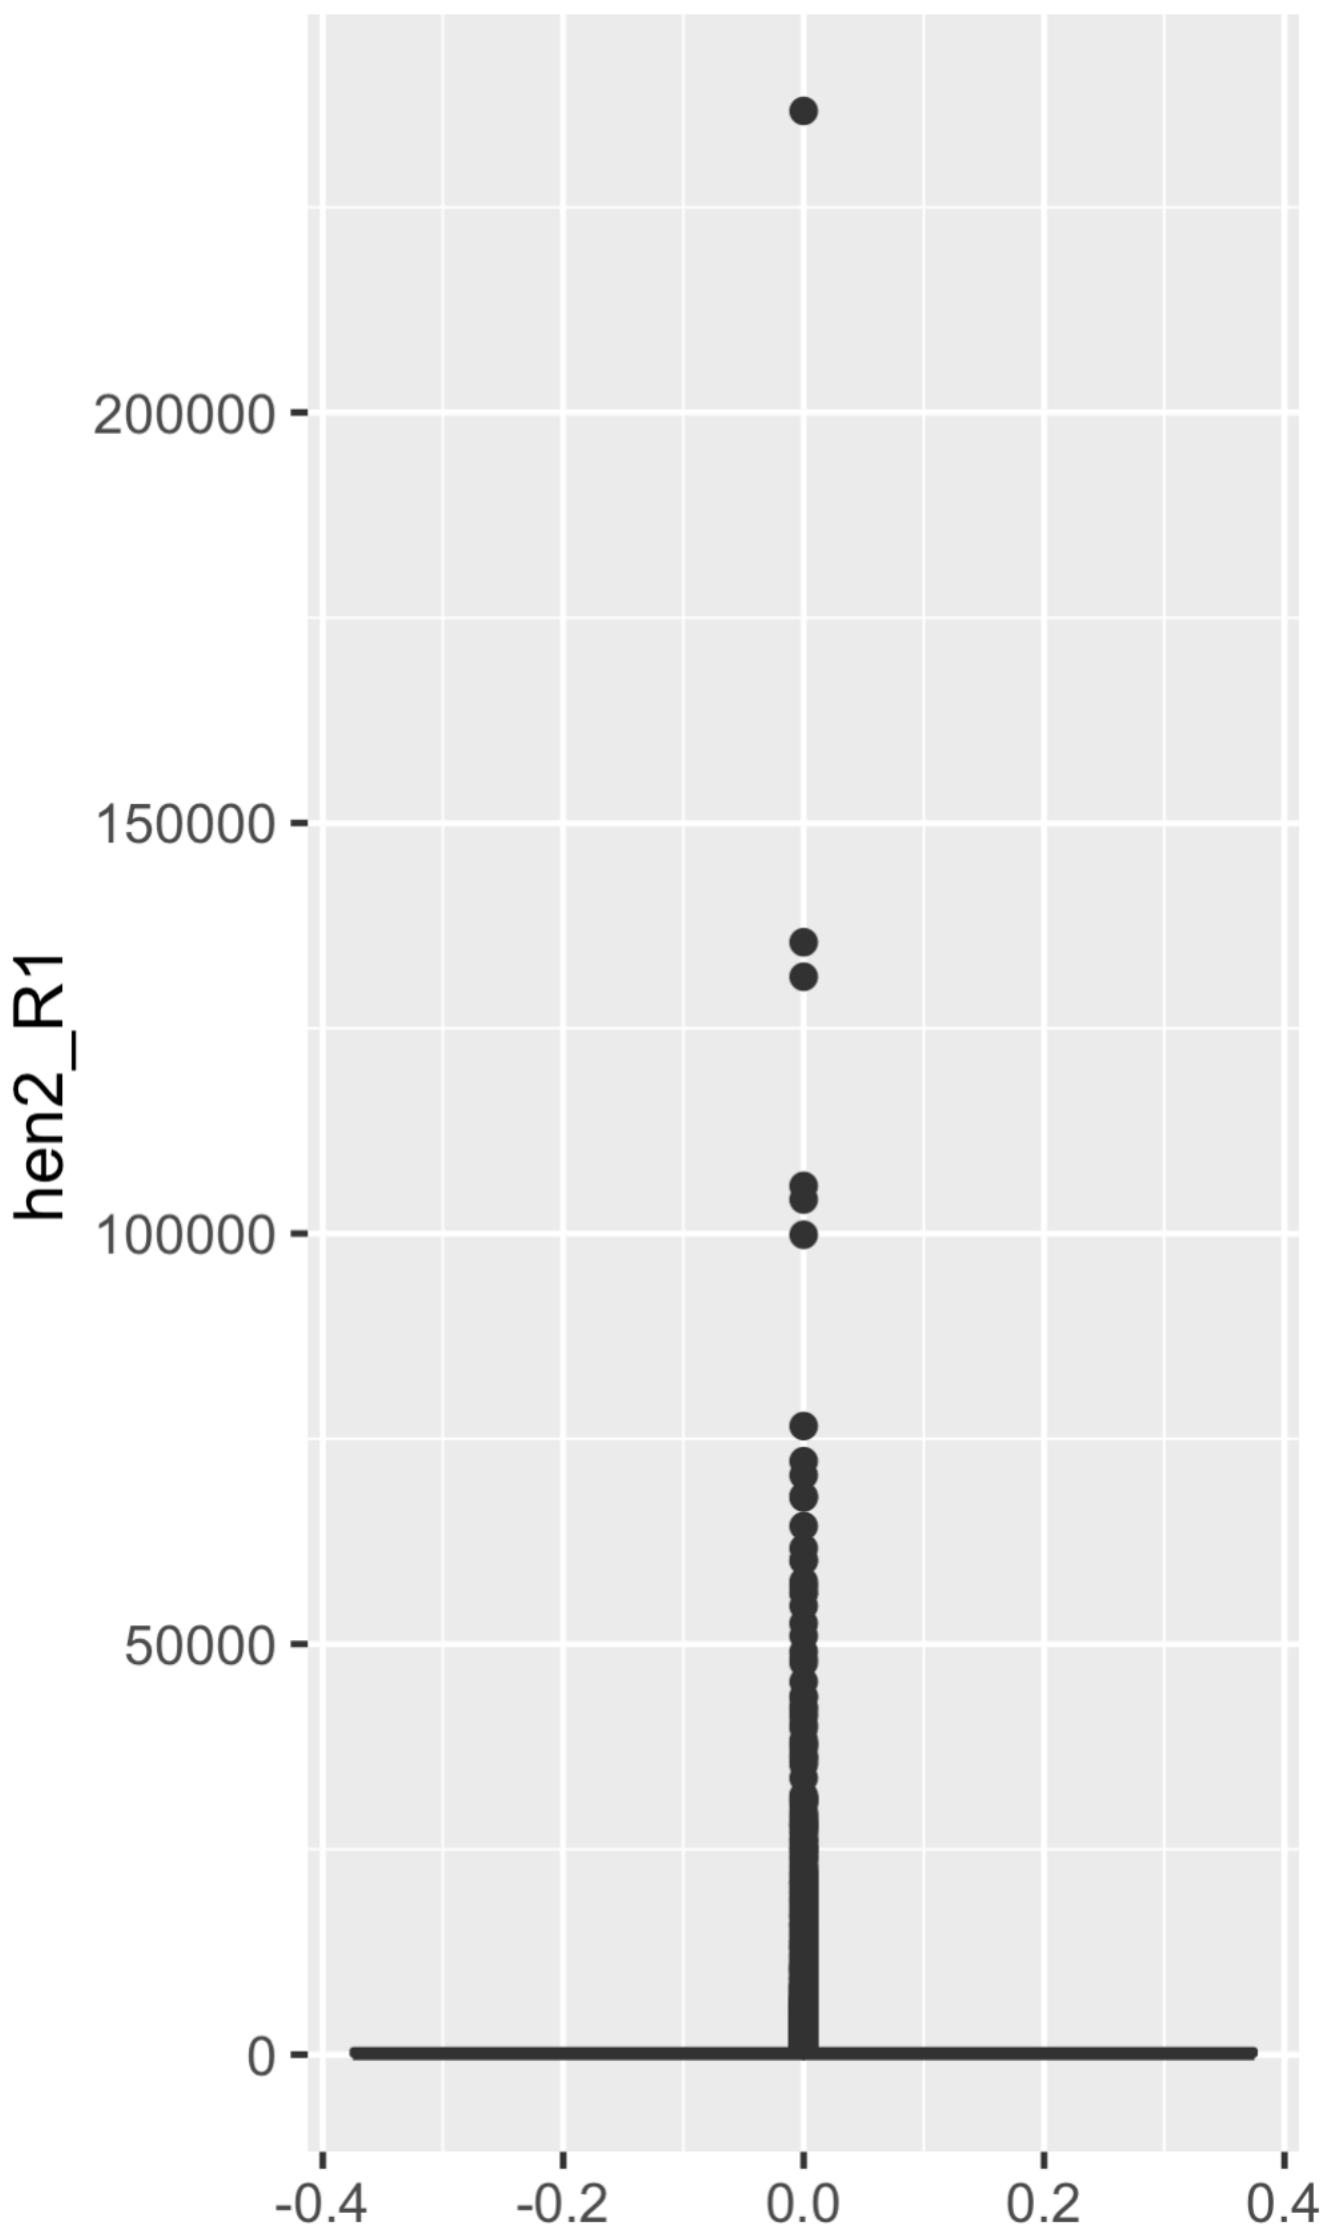
- What is happening?
- => Because some CAGE TCs have extremely high expression, they pull the scale up and crunch the distribution. We cannot see the box, but a lot of points (box plot “outliers”).
- Solution: log-scale!
/!\ Do not log-transform the **data**, log-transform the graphical scale!



Our first box plot

```
ggplot(data=TCs, aes(x='hen2 rep 1', y=hen2_R1)) +  
  geom_boxplot()
```

- What is happening?
- => Because some CAGE TCs have extremely high expression, they pull the scale up and crunch the distribution. We cannot see the box, but a lot of points (box plot “outliers”).
- Solution: log-scale!
/!\ Do not log-transform the **data**, log-transform the graphical scale!



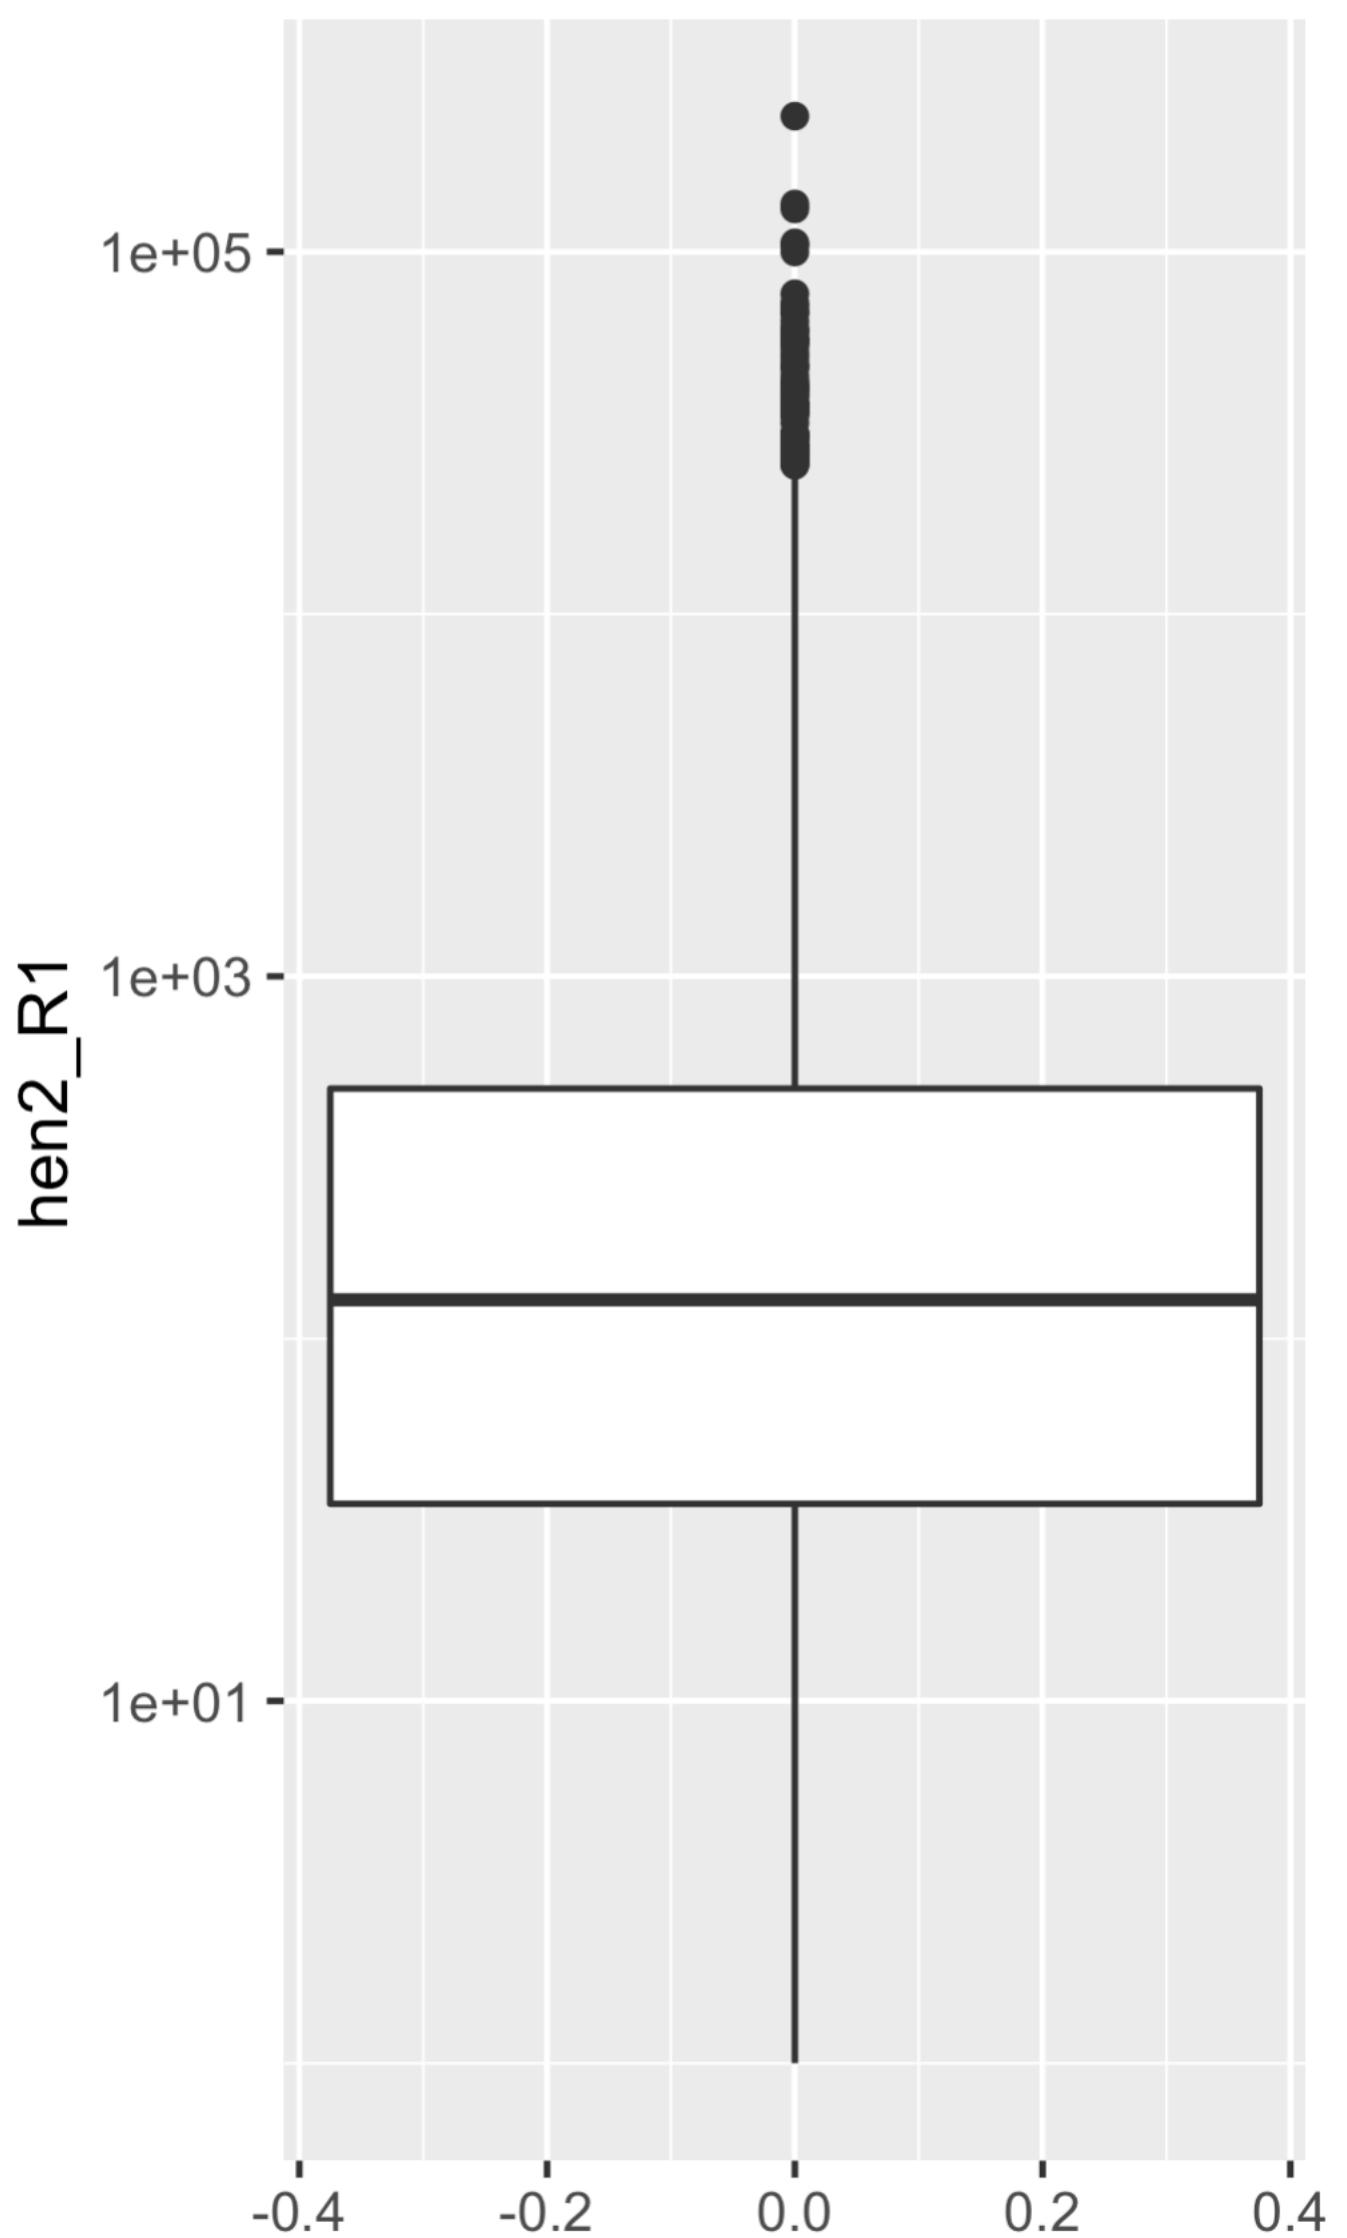
Our first box plot

```
ggplot(data=TCs, aes(x='hen2 rep 1', y=hen2_R1)) +  
  geom_boxplot() +  
  scale_y_log10()
```

- The plot is much better, but...
- Tidylog outputs some warning: some data points have been lost... why?!

```
Warning messages:  
1: Transformation introduced infinite values in continuous y-axis  
2: Removed 158 rows containing non-finite values (stat_boxplot).
```

- Tips: try `sum(TCs$hen2_R1 == 0)`



Comparing samples using box plots

- Problem: all of our coordinate parameters ($x=...$, $y=...$) are already occupied. We cannot indicate more samples (aka more columns) in that way.

```
ggplot(data=TCs, aes(x='hen2 rep 1', y=hen2_R1)) +  
  geom_boxplot() +  
  scale_y_log10()
```

- Melt the data frame 🔥

Wide & long format

- A typical data frame (or Excel sheet) is in the wide format
- Rows = variables (f.ex. samples)
Columns = observations (f.ex. genes or CAGE TCs)
Cells = values (f.ex. expression)
- It is called “**wide**” because such data frame is usually... wide.

country	year	cases	population
Afghanistan	1990	745	1908071
Afghanistan	2000	2666	2059360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	211258	127291272
China	2000	210766	128042583

variables

country	year	cases	population
Afghanistan	1990	745	1908071
Afghanistan	2000	2666	2059360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	211258	127291272
China	2000	210766	128042583

observations

country	year	cases	population
Afghanistan	1990	745	1908071
Afghanistan	2000	2666	2059360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	211258	127291272
China	2000	210766	128042583

values

Wide & long format

- A so-called “tidy” data frame will have a **long** format: each row is ONE value associated to one variable and one observation

WIDE

id	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
1	5.1	3.5	1.4	0.2
2	4.9	3.0	1.4	0.2
3	4.7	3.2	1.3	0.2
4	4.6	3.1	1.5	0.2
5	5.0	3.6	1.4	0.2

LONG

id	variable	value
1	Sepal.Length	5.1
2	Sepal.Length	4.9
3	Sepal.Length	4.7
4	Sepal.Length	4.6
5	Sepal.Length	5.0
1	Sepal.Width	3.5
2	Sepal.Width	3.0
3	Sepal.Width	3.2
4	Sepal.Width	3.1
5	Sepal.Width	3.6
1	Petal.Length	1.4
2	Petal.Length	1.4
3	Petal.Length	1.3
4	Petal.Length	1.5
5	Petal.Length	1.4
1	Petal.Width	0.2
2	Petal.Width	0.2
3	Petal.Width	0.2
4	Petal.Width	0.2
5	Petal.Width	0.2

Wide & long format

- A wide data frame allows to map each column (via the `aes()`) to a geometric layer of ggplot2
- Conceptually:
`ggplot(long_df, aes(x=variable, y=value) +
geom_boxplot())`
- Reformatting a data frame from **wide to long** format can be done with the `melt()` function from the `reshape2` package
- **Note:** The opposite, from **long to wide**, can be done with `pivot_wider()`
- Let's try it with out CAGE TC expression boxplot

LONG		
id	variable	value
1	Sepal.Length	5.1
	Sepal.Length	4.9
	Sepal.Length	4.7
	Sepal.Length	4.6
	Sepal.Length	5.0
2	Sepal.Width	3.5
	Sepal.Width	3.0
	Sepal.Width	3.2
	Sepal.Width	3.1
	Sepal.Width	3.6
3	Petal.Length	1.4
	Petal.Length	1.4
	Petal.Length	1.3
	Petal.Length	1.5
	Petal.Length	1.4
4	Petal.Width	0.2
	Petal.Width	0.2

Wide & long format

- A wide data frame allows to map each column (via the `aes()`) to a geometric layer of ggplot2
- Conceptually:
`ggplot(long_df, aes(x=variable, y=value)) +
geom_boxplot()`
- Reformatting a data frame from wide to long format can be done with the `melt()` function from the reshape2 package
- Let's try it with our CAGE TC expression boxplot

Wide & long format

- Reproduce the following command and explain what happens

```
12 # wide format
13 TCs
14 # melt df to long format, using TC_id as "key"
15 TCs %>%
16   melt(id.vars='TC_id') %>%
17   head()
```

- **Note:** our TCs data frame used to be a tibble. Certain operations on tibbles coerce a tibble back to a data frame. Use of the head() command allows to not be flooded by the whole data frame content (à la UNIX “cat”).

Wide & long format

- Let's jump ahead:

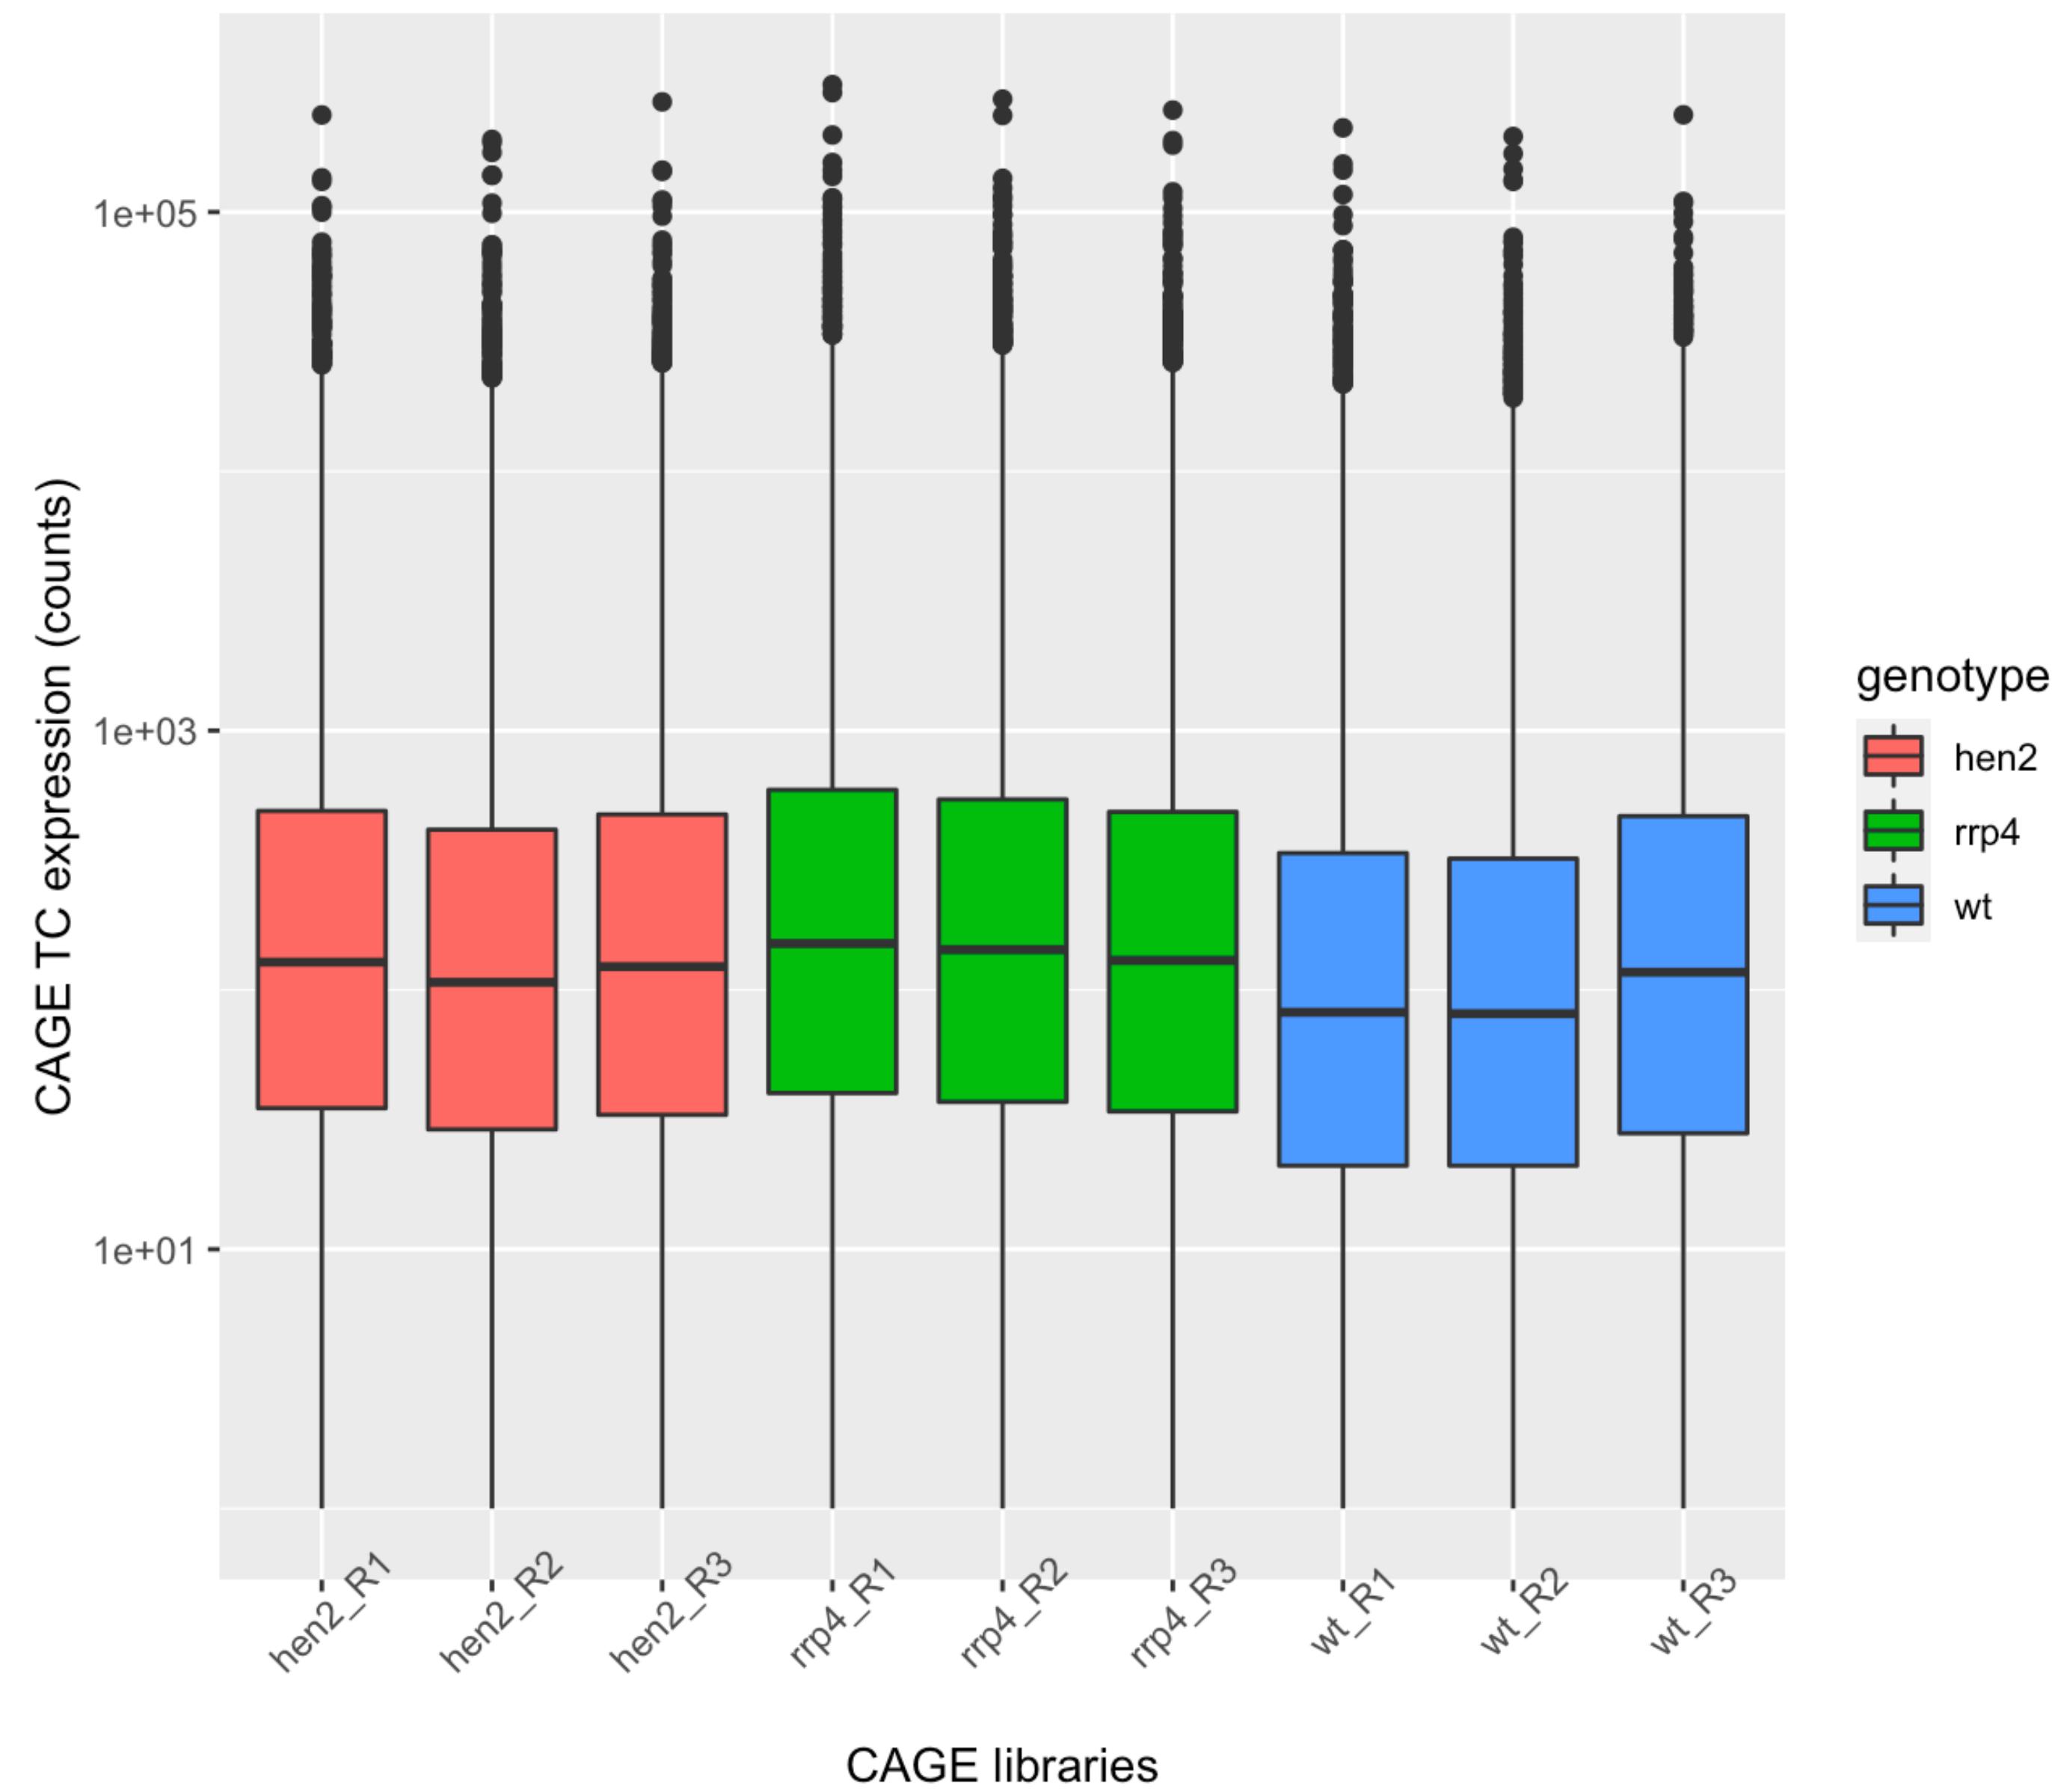
Reproduce the code **line by line**
and dissect what happens

```
# wide format  
TCs  
# melt df to long format, using TC_id as "key", and plot  
TCs %>%  
  melt(id.vars='TC_id') %>%  
  separate(col='variable', into=c('genotype', 'replicate'), sep='_', remove=F) %>%  
  ggplot(aes(x=variable, y=value, fill=genotype)) +  
    geom_boxplot() +  
    scale_y_log10() +  
    theme(axis.text.x=element_text(angle=45),  
          aspect.ratio=1) +  
    labs(x='CAGE libraries', y='CAGE TC expression (counts)',  
         title='This starts to look like something',  
         subtitle='there is even place for a subtitle!',  
         caption='Caption: Raw tag counts')
```

Wide & long format

- RESULT:

This starts to look like something
there is even place for a subtitle!



Caption: Raw tag counts

Step 1: go from wide to long format

Wide & long format

- Let's jump ahead:

```
# wide format  
TCs  
# melt df to long format, using TC_id as "key", and plot  
TCs %>%  
  melt(id.vars='TC_id')
```

- Here we use the column “TC_id” as the “key”

Wide & long format

- Let's jump ahead:

```
# wide format  
TCs  
# melt df to long format, using TC_id as "key", and plot  
TCs %>%  
  melt(id.vars='TC_id') %>%  
  separate(col='variable', into=c('genotype', 'replicate'), sep='_', remove=F)
```

- **separate()** is a function that splits a column into several columns, based on a character string (can be a regex!)
- one option is to not have the original column deleted (remove=F)

Step 2: extract genotype information only (not replicate) by splitting a column based on underscore

Wide & long format

- Let's jump ahead:

```
# wide format  
TCs  
# melt df to long format, using TC_id as "key", and plot  
TCs %>%  
  melt(id.vars='TC_id') %>%  
  separate(col='variable', into=c('genotype', 'replicate'), sep='_', remove=F) %>%  
  ggplot(aes(x=variable, y=value, fill=genotype))
```

- We can now initiate the ggplot() call, and use one column for the X variable, one for the values (Y), and use a filling color (fill=) thanks to the genotype column
- Important: ggplot2 will use different fills for each genotype, and hence will create one boxplot for each genotype automatically, this is the power of the grammar of graphics

Step 3: pass to ggplot() and map aesthetic (aes) to columns

Step 4: box plot on a log-scale

Wide & long format

- Let's jump ahead:

```
# wide format  
TCs  
# melt df to long format, using TC_id as "key", and plot  
TCs %>%  
  melt(id.vars='TC_id') %>%  
    separate(col='variable', into=c('genotype', 'replicate'), sep='_', remove=F) %>%  
    ggplot(aes(x=variable, y=value, fill=genotype)) +  
      geom_boxplot() +  
      scale_y_log10()
```

- We impose a couple of supplemental graphic layers (+) that will specify a boxplot and a log10-scaled Y-axis

Wide & long format

- Let's jump ahead:

```
# wide format  
TCs  
# melt df to long format, using TC_id as "key", and plot  
TCs %>%  
  melt(id.vars='TC_id') %>%  
  separate(col='variable', into=c('genotype', 'replicate'), sep='_', remove=F) %>%  
  ggplot(aes(x=variable, y=value, fill=genotype)) +  
    geom_boxplot() +  
    scale_y_log10() +  
    theme(axis.text.x=element_text(angle=45),  
          aspect.ratio=1)
```

Step 5: Improve the plot's general parameters, including layout

- The theme() layer controls a lot of theme-related features, such as the axis labels, overall plot aspect ratio, and SO many more.

Wide & long format

- Let's jump ahead:

```
# wide format  
TCs  
# melt df to long format, using TC_id as "key", and plot  
TCs %>%  
  melt(id.vars='TC_id') %>%  
  separate(col='variable', into=c('genotype', 'replicate'), sep='_', remove=F) %>%  
  ggplot(aes(x=variable, y=value, fill=genotype)) +  
    geom_boxplot() +  
    scale_y_log10() +  
    theme(axis.text.x=element_text(angle=45),  
          aspect.ratio=1)  
  labs(x='CAGE libraries', y='CAGE TC expression (counts)',  
       title='This starts to look like something',  
       subtitle='there is even place for a subtitle!',  
       caption='Caption: Raw tag counts')
```

Step 6: Make the plot understandable
with legends (labels: `labs()`)!!!
=> very important.

But... why?

- In this block of code, you:
 - 1. Reshaped the data (data wrangling)
 - 2. Extracted genotype information for each expression value (separate)
 - 3. Summarised the expression distribution of each sample
(boxplot with quartiles, outliers, etc)
 - 4. Visualized the data with genotype indication (fill=genotype)
 - 5. Scaled the visualisation to make the figure intelligible
 - 6. Set a plot layout with all the essential scientific info (labels)
- In one command, very close to simple English
- Without altering the underlying dataset (TCs data frame)
- Without creating temporary variables

Normalize expression: Raw counts vs TPM

- **TPM** = Tags Per Million mapped tags
- TPM = normalisation of expression data with regards to sequencing depth
 - accounts for technical sequencing differences
 - doesn't account for biological differences as RPKM
- **RPKM** = Reads Per Kilobase per Million mapped reads
 - accounts for technical sequencing differences
 - accounts for each gene's length (longer gene = more reads)
- RPKM is typically needed for RNA-seq when we need to compare different genes, but not when comparing the same gene between samples
- Since CAGE only captures TSSs, read length does not weigh in, so TPM is enough

Read more on this here: www.reneshbedre.com/blog/expression_units.html

Let's normalise our CAGE TC expression

Difficult exercise

- Formula: $TPM = \text{tags count} / (\text{total mapped tags} / 1,000,000)$
- **1a)** Compute the sum of reads (aka tags) for each sample
1b) and divide it by 1,000,000
1c) and save it as “`TPM_factor`”

Let's normalise our CAGE TC expression

Difficult exercise

- Formula: $TPM = \text{tags count} / (\text{total mapped tags} / 1,000,000)$
- Apply the formula above:
 - 2a) Divide each gene's tag count by the TPM factor of the relevant library.
 - Tips 1: this is dividing each ROW of the TCs data frame by the vector of TPM_factor, the sort of operation `apply()` is good at ;)
 - Tips 2: look at the `dim()` before and after `apply()`, what happened? could `t()` be useful? What does `t()` do?
 - 2b) Save the TPM-normalized data frame as “TCs TPM”
 - 2c) Does the sum of normalised TPM per sample amounts to 1million?

Let's normalise our CAGE TC expression

Difficult exercise

- Formula: $TPM = \text{tags count} / (\text{total mapped tags} / 1,000,000)$
- 3) Using our new data frame:
Make a boxplot of the normalised CAGE TC expression across samples, color-filled by genotype
- 4) What's different compared to the raw count box plots?